


SIG Cluster Lifecycle

Overview and projects

Martin Vladev - SAP

Rostislav M. Georgiev - VMware



SIGs?



“Special Interest Groups (SIGs) are persistent open groups that focus on a part of the project.”

--- Kubernetes Community README

SIGs?



API Machinery

Azure

Docs

GCP

Architecture

Big Data

IBM Cloud

Instrumentation

Apps

CLI

Cloud Provider

Multicluster

Auth

Cluster Lifecycle

Network

Node

Autoscaling

Contributor Experience

Cluster Ops

AWS

OpenStack

PM

Release

Scalability

Scheduling

Service Catalog

Storage

Testing

UI

VMware

Windows



kubernetes


SIG Cluster Lifecycle



“SIG Cluster Lifecycle’s objective is to simplify creation, configuration, upgrade, downgrade, and teardown of Kubernetes clusters and their components.”


--- SIG Cluster Lifecycle Charter

SIG Cluster Lifecycle



kubernetes

SIG Cluster Lifecycle



kubernetes



KUBEADM

Kubeadm Design Goals



- Different deployment options
- Public & private clouds, VMs & bare metal
- Different Linux distros
- Different CRIs
- Flexible network options

Kubeadm Design Goals



- Upgrades
- Easy workflow
 - Provision one machine at a time
 - End users can use it
 - Automated deployment tools

Prerequisites



Machine 1

- Operating System

Machine 2

- Operating System

Machine 3

- Operating System



Prerequisites



Machine 1

- Container Runtime
- Operating System

Machine 2

- Container Runtime
- Operating System

Machine 3

- Container Runtime
- Operating System



Prerequisites



Machine 1

- Kubelet + kubeadm
- Container Runtime
- Operating System

Machine 2

- Kubelet + kubeadm
- Container Runtime
- Operating System

Machine 3

- Kubelet + kubeadm
- Container Runtime
- Operating System



Prerequisites



Control Plane

- Kubelet + kubeadm
- Container Runtime
- Operating System

Worker 1

- Kubelet + kubeadm
- Container Runtime
- Operating System

Worker 2

- Kubelet + kubeadm
- Container Runtime
- Operating System



kubeadm init



```
rosti@ubuntu:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] using Kubernetes version: v1.12.1
[preflight] running pre-flight checks
[preflight/images] Pulling images required for setting up a Kubernetes cluster
[preflight/images] This might take a minute or two, depending on the speed of your internet connection
[preflight/images] You can also perform this action in beforehand using 'kubeadm config images pull'
```



Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

You can now join any number of machines by running the following on each node
as root:

```
kubeadm join 192.168.217.129:6443 --token g7qg8j.vj4blgjdsagr9x0 --discovery-token-ca-cert-hash sha256:ddf6a76a90edbd47ddbeb3fea4a06180d1f0b26eea9f91db42f207baadc2b43f
```



kubernetes

Kubeadm init



```
rosti@ubuntu:~$ kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-576cbf47c7-glxfj9	0/1	ContainerCreating	0	72s
kube-system	coredns-576cbf47c7-zxp7r	0/1	ContainerCreating	0	72s
kube-system	etcd-ubuntu	1/1	Running	0	9s
kube-system	kube-apiserver-ubuntu	1/1	Running	0	24s
kube-system	kube-controller-manager-ubuntu	1/1	Running	0	37s
kube-system	kube-proxy-b796d	1/1	Running	0	72s
kube-system	kube-scheduler-ubuntu	1/1	Running	0	16s



kubernetes

kubeadm init



Control Plane

- PODs
- Kubelet + kubeadm
- Container Runtime
- Operating System

☐ API Server

☐ Controller Manager

☐ Scheduler

☐ Kube-Proxy

☐ etcd

☐ CoreDNS

Kubelet + Config



kubernetes

kubeadm init



```
rosti@ubuntu:~$ kubectl apply -f https://raw.githubusercontent.com/clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.extensions/kube-flannel-ds-amd64 created
daemonset.extensions/kube-flannel-ds-arm64 created
daemonset.extensions/kube-flannel-ds-arm created
daemonset.extensions/kube-flannel-ds-ppc64le created
daemonset.extensions/kube-flannel-ds-s390x created
```



kubernetes

kubeadm init



```
rosti@ubuntu:~$ kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-576cbf47c7-glxfj9	1/1	Running	0	4m20s
kube-system	coredns-576cbf47c7-zxp7r	1/1	Running	0	4m20s
kube-system	etcd-ubuntu	1/1	Running	0	3m17s
kube-system	kube-apiserver-ubuntu	1/1	Running	0	3m32s
kube-system	kube-controller-manager-ubuntu	1/1	Running	0	3m45s
kube-system	kube-flannel-ds-amd64-c6kjt	1/1	Running	0	2m56s
kube-system	kube-proxy-b796d	1/1	Running	0	4m20s
kube-system	kube-scheduler-ubuntu	1/1	Running	0	3m24s



kubernetes

kubeadm init



Control Plane

- PODs
- Kubelet + kubeadm
- Container Runtime
- Operating System



API Server



Controller Manager



Scheduler



Kube-Proxy



etcd



CoreDNS



Flannel

Kubelet + Config



kubernetes

kubeadm join



```
rosti@worker1:~$ sudo kubeadm join 192.168.217.129:6443 --token g7qg8j.vj4blgjdbdsagr9x0 --discovery-token-ca-cert-hash sha256:ddf6a76a90edbd47ddbeb3fea4a06180d1f0b26eea9f91db42f207baadc2b43f
[preflight] running pre-flight checks
[discovery] Trying to connect to API Server "192.168.217.129:6443"
[discovery] Created cluster-info discovery client, requesting info from "https://192.168.217.129:6443"
[discovery] Requesting info from "https://192.168.217.129:6443" again to validate TLS against the pinned public key
[discovery] Cluster info signature and contents are valid and TLS certificate validates against pinned roots, will use API Server "192.168.217.129:6443"
[discovery] Successfully established connection with API Server "192.168.217.129:6443"
[kubelet] Downloading configuration for the kubelet from the "kubelet-config-1.12" ConfigMap in the kube-system namespace
[kubelet] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[preflight] Activating the kubelet service
[tlsbootstrap] Waiting for the kubelet to perform the TLS Bootstrap...
[patchnode] Uploading the CRI Socket information "/var/run/dockershim.sock" to the Node API object "worker1" as an annotation

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the master to see this node join the cluster.

rosti@worker1:~$
```



kubernetes

kubeadm join



```
rosti@ubuntu:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ubuntu	Ready	master	16m	v1.12.1
worker1	Ready	<none>	72s	v1.12.1

```
rosti@ubuntu:~$ kubectl get pods --all-namespaces -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE
kube-system	coredns-576cbf47c7-glxfj9	1/1	Running	0	46m	10.244.0.9	ubuntu	<none>
kube-system	coredns-576cbf47c7-zxp7r	1/1	Running	0	46m	10.244.0.10	ubuntu	<none>
kube-system	etcd-ubuntu	1/1	Running	0	33m	192.168.217.129	ubuntu	<none>
kube-system	kube-apiserver-ubuntu	1/1	Running	0	33m	192.168.217.129	ubuntu	<none>
kube-system	kube-controller-manager-ubuntu	1/1	Running	0	33m	192.168.217.129	ubuntu	<none>
kube-system	kube-flannel-ds-amd64-6zzsd	1/1	Running	0	31m	192.168.217.130	worker1	<none>
kube-system	kube-flannel-ds-amd64-c6kjt	1/1	Running	0	44m	192.168.217.129	ubuntu	<none>
kube-system	kube-proxy-4kmgk	1/1	Running	0	31m	192.168.217.130	worker1	<none>
kube-system	kube-proxy-b796d	1/1	Running	0	46m	192.168.217.129	ubuntu	<none>
kube-system	kube-scheduler-ubuntu	1/1	Running	0	33m	192.168.217.129	ubuntu	<none>



kubeadm



- The Future
 - Going GA in 1.13
 - Used by Cluster API
- Feedback
 - Have you considered kubeadm?
 - What is missing?
 - What is done wrong?
 - Can we do things better?

kubeadm



- Join the kubeadm office hours meetings
 - Wednesdays, 7PM BG
 - [Meeting notes and agenda](#)
- Ping us on Slack
 - #sig-cluster-lifecycle
 - #kubeadm



kubernetes

kubeadm



kubernetes / kubeadm

Watch 84

Star 709

Fork 126

Code

Issues 111

Pull requests 2

Projects 0

Wiki

Insights

Filters

is:issue is:open

Labels

Milestones

New issue

111 Open 983 Closed

Author

Labels

Projects

Milestones

Assignee

Sort

i am unable to get all my pods in running state priority/awaiting-more-evidence

13

#1179 opened 3 hours ago by MohanNagendraKumar

Bad reference to controlManagerExtraArgs in v1alpha3 documentation good first issue help wanted

5

kind/documentation lifecycle/active

#1178 opened 14 hours ago by nkinkade

Refactor directory fetching code good first issue kind/refactor priority/important-longterm

#1177 opened 17 hours ago by chuckha

During Kubelet startup on node reboot, docker cri networking is used area/ecosystem help wanted

#1176 opened 17 hours ago by ChiefAlexander

Disabling CoreDNS featuregate does not disable postupgrade task help wanted

#1175 opened 20 hours ago by mattymo



kubernetes



Cluster API

Overview



- Declarative, Kubernetes-style APIs to cluster creation, configuration, and management
- Extensible and environment-agnostic API
- Ease and standardize Node management and operation
 - Cluster autoscaler
 - Node Problem detector

<https://sigs.k8s.io/cluster-api>



kubernetes

API Resources 1/2



- Cluster
 - Describes a single cluster
 - Pod, Service and Domain
 - Provider-specific configuration
- Machine
 - single Node
 - Kubelet and Control plane configuration
 - Provider-specific configuration

API Resources 2/2



- MachineSet
 - Creates multiple Machine replicas
 - Same behavior as ReplicaSet
- MachineDeployment
 - Manages multiple ReplicaSet
 - Same behavior as Deployment
 - Rolling updates of Machines



```
apiVersion: cluster.k8s.io/v1alpha1
kind: Cluster
metadata:
  name: test1-ckf8e
spec:
  clusterNetwork:
    services:
      cidrBlocks: ["10.96.0.0/12"]
    pods:
      cidrBlocks: ["192.168.0.0/16"]
      serviceDomain: cluster.local
  providerConfig:
    value:
      apiVersion: gceproviderconfig/v1alpha1
      kind: GCEClusterProviderConfig
      project: some-gcp-project
```



```
apiVersion: cluster.k8s.io/v1alpha1
kind: Cluster
metadata:
  name: test1-ckf8e
spec:
  clusterNetwork:
    services:
      cidrBlocks: ["10.96.0.0/12"]
    pods:
      cidrBlocks: ["192.168.0.0/16"]
    serviceDomain: cluster.local
  providerConfig:
    value:
      apiVersion: gceproviderconfig/v1alpha1
      kind: GCEClusterProviderConfig
      project: some-gcp-project
```

Provider-specific config



```
apiVersion: cluster.k8s.io/v1alpha1
kind: Machine
metadata:
  generateName: gce-master-
  labels:
    set: master
spec:
  providerConfig:
    value:
      apiVersion: gceproviderconfig/v1alpha1
      kind: GCEMachineProviderConfig
      roles:
        - Master
      zone: europe-west2-a
      machineType: n1-standard-2
      os: ubuntu-1604-lts
      disks:
        - initializeParams:
            diskSizeGb: 30
            diskType: pd-standard
  versions:
    kubelet: 1.12.0
    controlPlane: 1.12.0
```



```
apiVersion: cluster.k8s.io/v1alpha1
kind: Machine
metadata:
  generateName: gce-master-
  labels:
    set: master
spec:
  providerConfig:
    value:
      apiVersion: gceproviderconfig/v1alpha1
      kind: GCEMachineProviderConfig
      roles:
        - Master
      zone: europe-west2-a
      machineType: n1-standard-2
      os: ubuntu-1604-lts
      disks:
        - initializeParams:
            diskSizeGb: 30
            diskType: pd-standard
  versions:
    kubelet: 1.12.0
    controlPlane: 1.12.0
```

Provider-specific config



Infrastructure Providers

Cloud, On-Prem, Bare Metal
Potential Providers =>



AWS



Azure



GCP

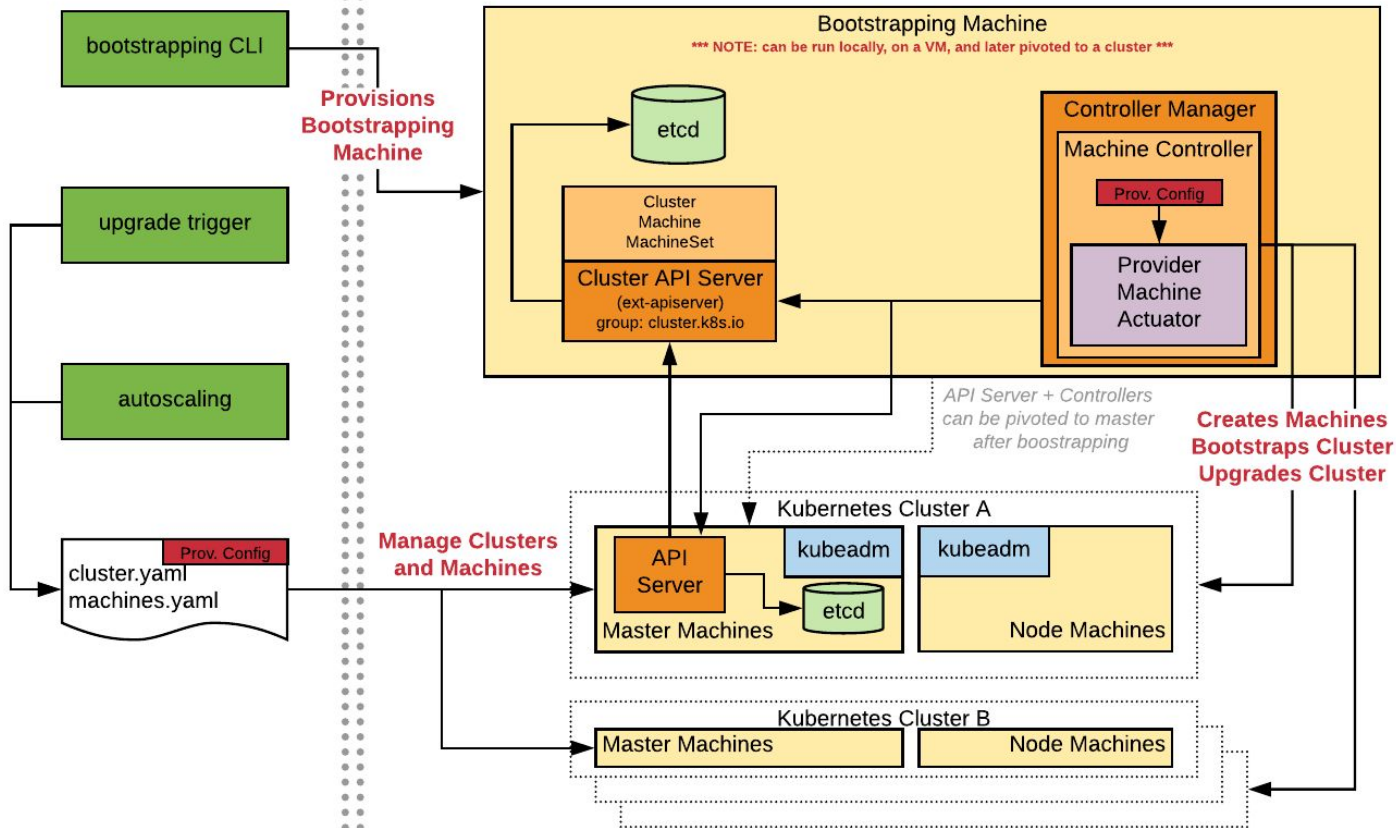


vmware



Bare Metal

 Provider Configuration



Thank you!



kubernetes