



# Introduction to service mesh



**kubernetes**



**Istio**

by Plamen Stoyanov

**[ K8S & CNCF Bulgaria] 25/10/2018**




# What is kubernetes good for

- Control and automate application deployments and updates.
- Make better use of hardware to maximize resources needed to run your apps.
- Scale containerized applications and their resources on the fly.
- Mount and add storage to run stateful apps.
- Declaratively manage services, which guarantees the deployed applications are always running how you deployed them.
- Health-check and self-heal your apps with autoplacement, autorestart, autoreplication, and autoscaling.

# k8s cluster

Kubernetes cluster consist of

- Master
- Worker nodes




# k8s node

A node is a worker machine in  
Kubernetes  
May be a VM or physical machine  
Multiple Pods can run on one Node

Every Kubernetes Node runs at least:


- Kubelet
- A container runtime



# k8s pod

A Pod is a group of one or more application containers

- Shared context
- Shared storage
- Unique cluster IP
- Tightly coupled
- Always run on the same node




# k8s Service


A Kubernetes **Service** is an abstraction layer which defines a logical set of Pods and enables external traffic exposure, load balancing and service discovery for those Pods.

Service:

- Exposes Pods to external traffic
- Load balancing traffic across multiple Pods
- Using labels



# Monolith to microservices









## Microservices challenges

- Observability
  - monitoring
  - performance
  - logging
- Traceability
- Security
  - authentication
  - encryption
  - auditing



**What is a service mesh ?**

Network for services, not for bytes



- Connect - resiliency, discovery, load balancing
- Manage - traffic control, policy enforcement
- Monitor - metrics, logging, tracing
- Secure - end-to-end authentication and authorization
- Deploy - canary, rolling upgrades, separation
- Test - manage chaos

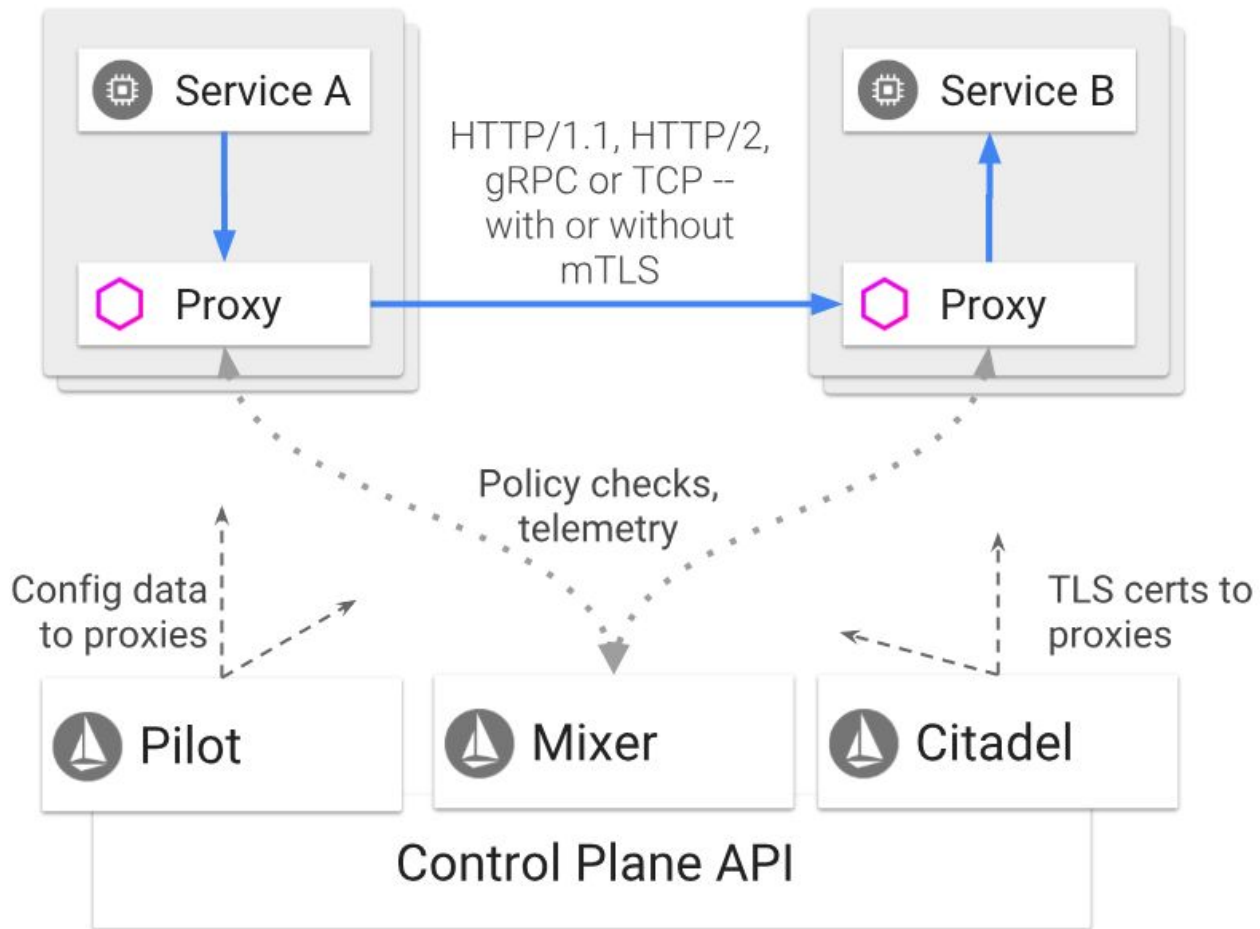


## Main Istio contributors



# Istio Architecture

- **Data plane** - composed of a set of intelligent proxies (Envoy) deployed as sidecars
- **Control plane** - manages and configures the proxies to route traffic



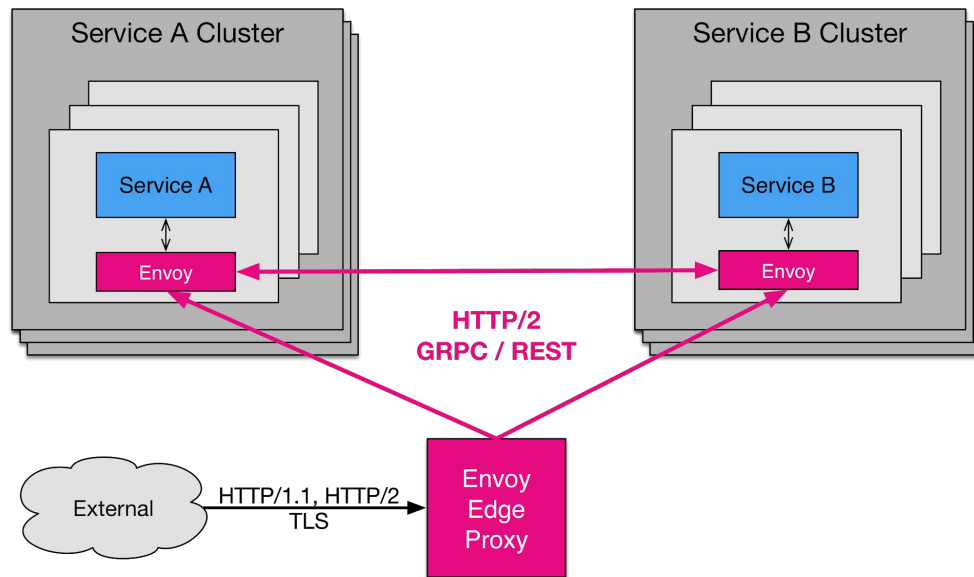
# Envoy

Istio uses extended version of the [Envoy](#) proxy.

Features:

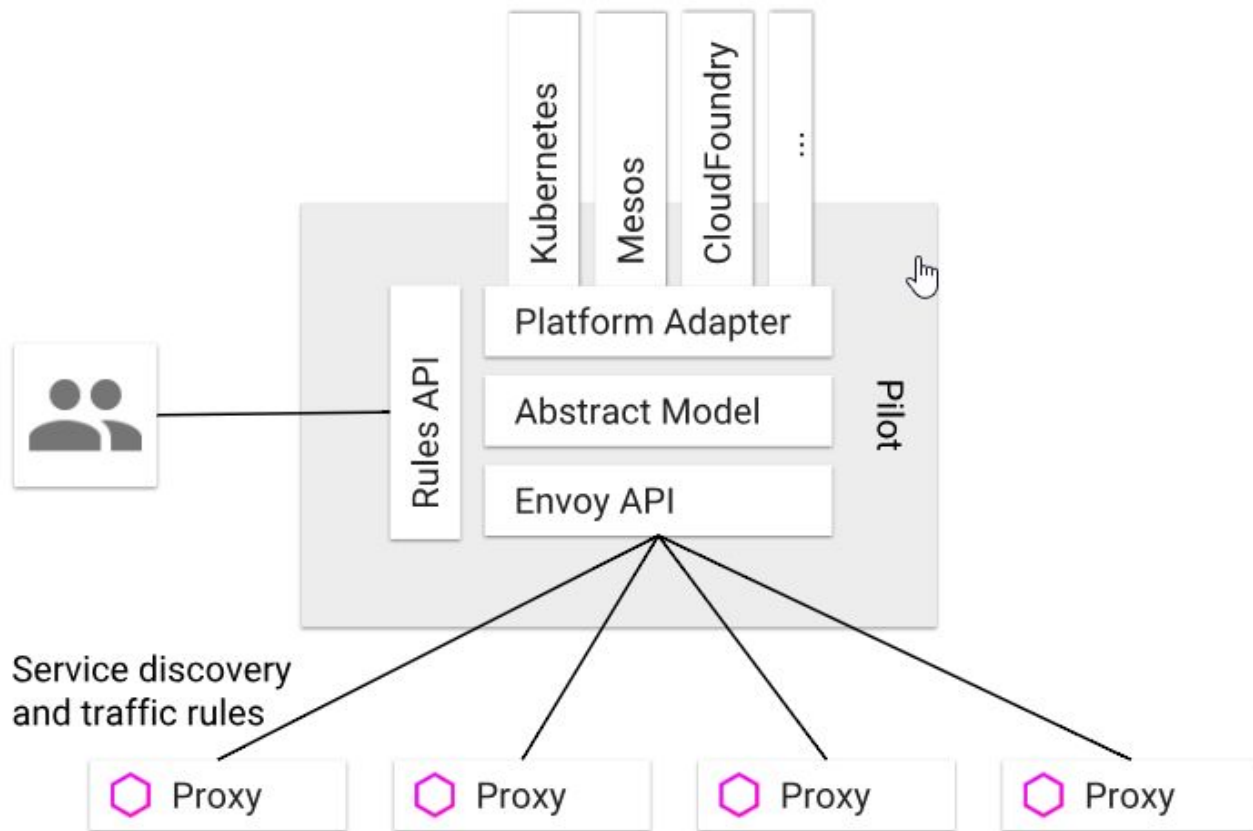
- Dynamic service discovery
- Load balancing
- TLS termination
- HTTP/2 and gRPC proxies
- Circuit breakers
- Health checks
- Staged rollouts with %-based traffic split
- Fault injection
- Rich metrics

Envoy is deployed as a **sidecar** to the relevant service in the same Kubernetes pod



# Pilot

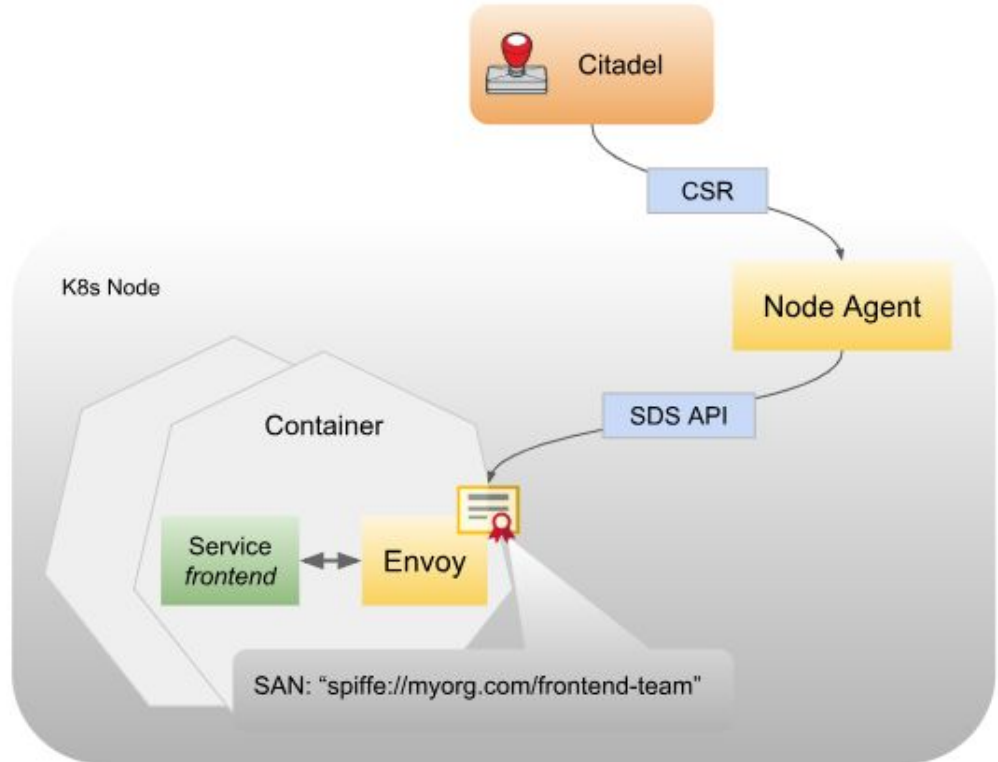
Pilot converts high level routing rules that control traffic behavior into Envoy-specific configurations, and propagates them to the sidecars at runtime



# Citadel

Citadel provides strong service-to-service and end-user authentication with built-in identity and credential management.

- SPIFFE
- PKI

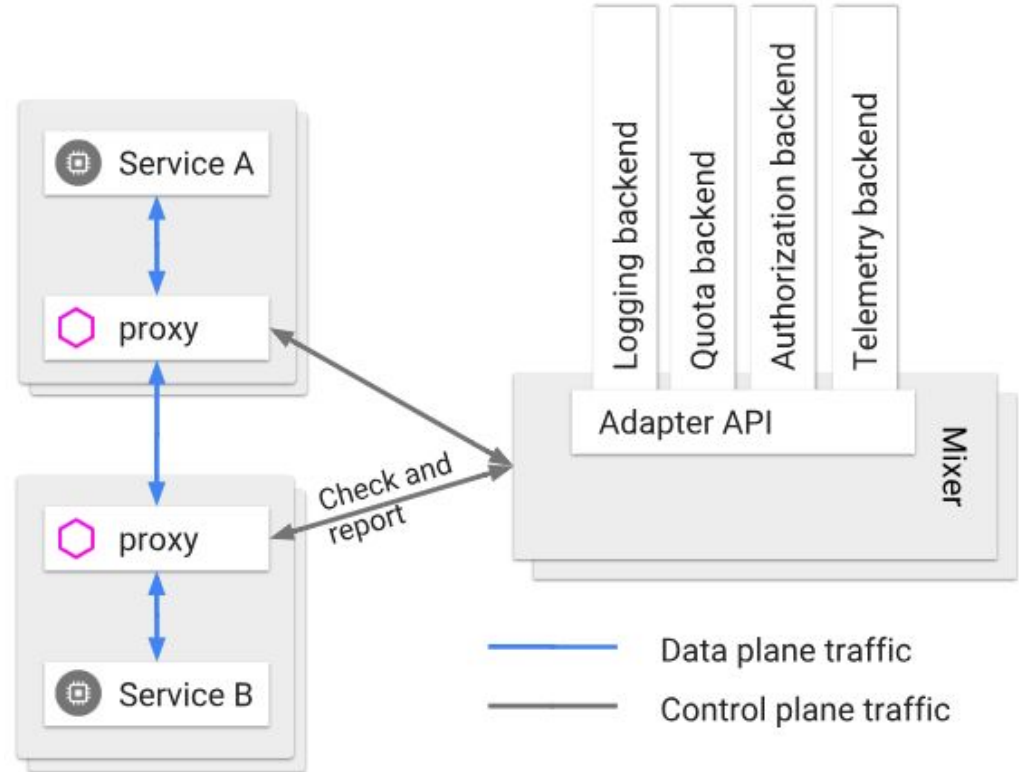




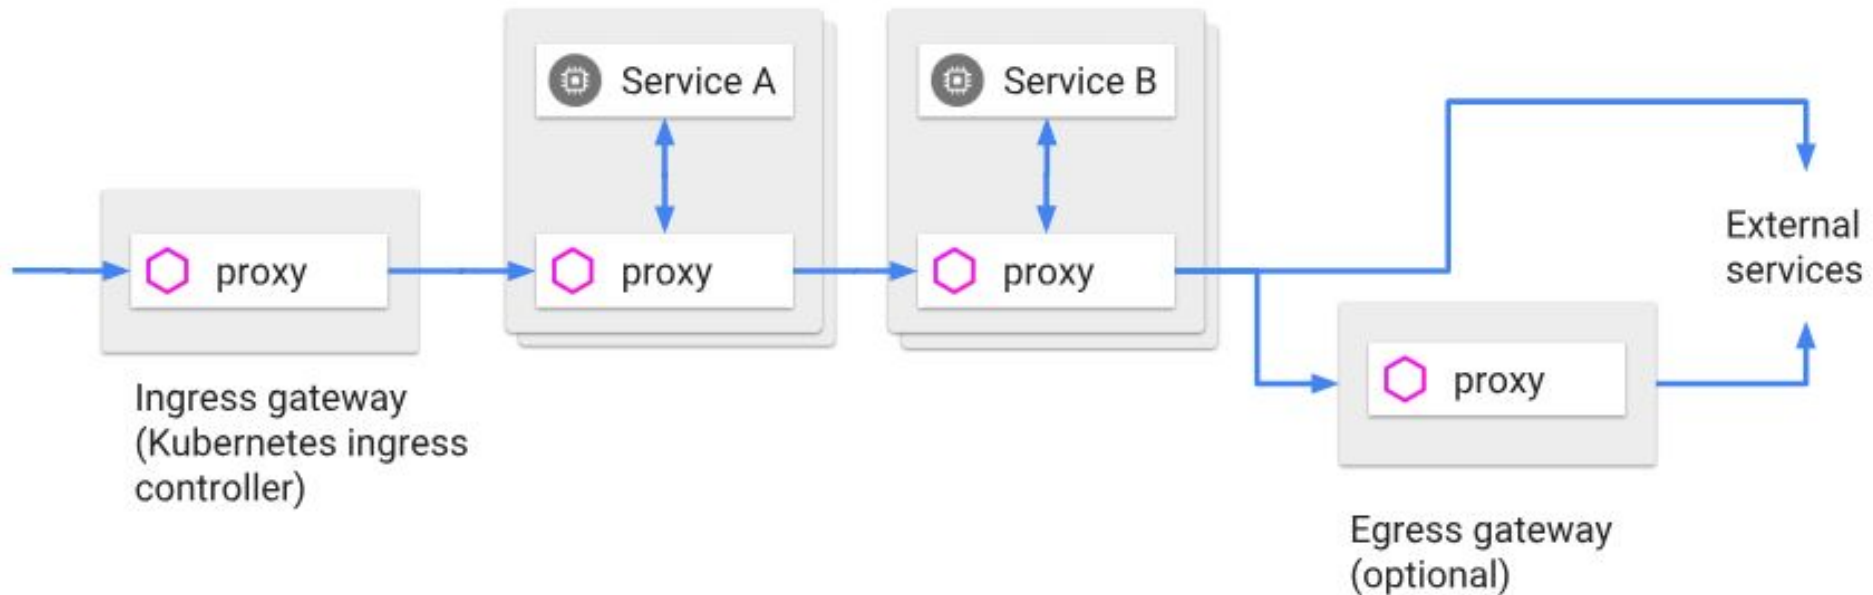
# Mixer

[Mixer](#) is a platform-independent component. Mixer enforces access control and usage policies across the service mesh, and collects telemetry data from the Envoy proxy and other services

Mixer includes a flexible plugin model.

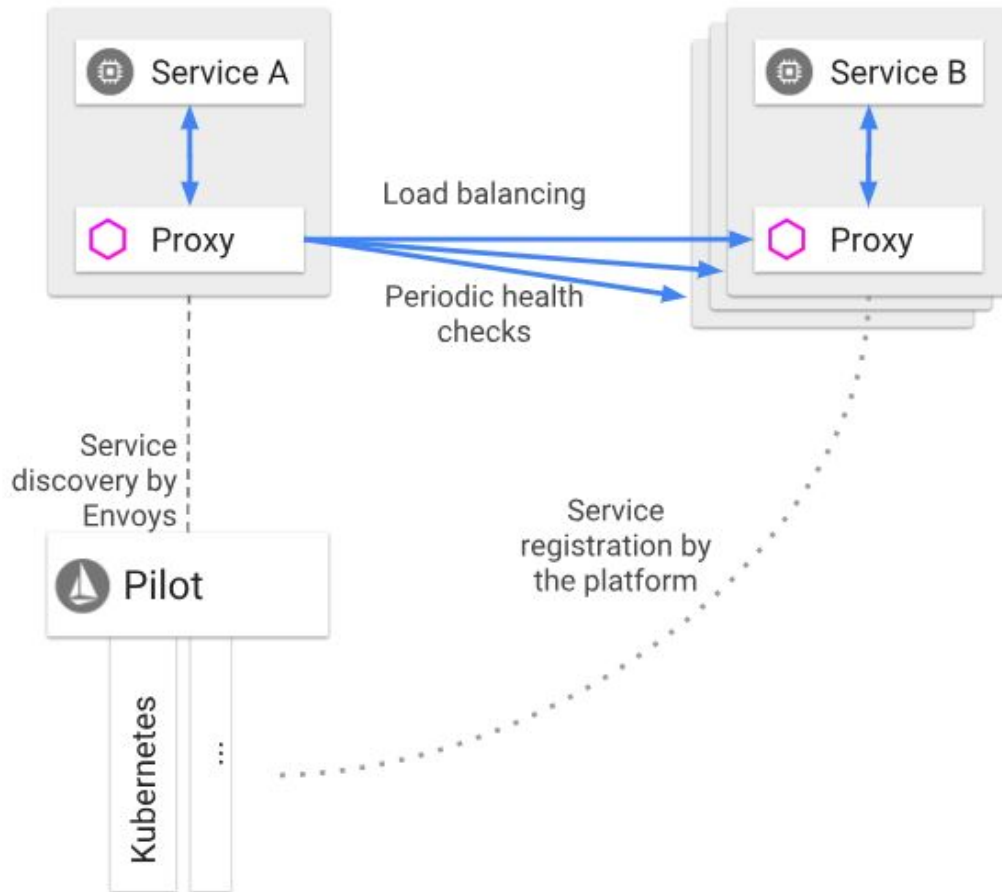


# Request flow



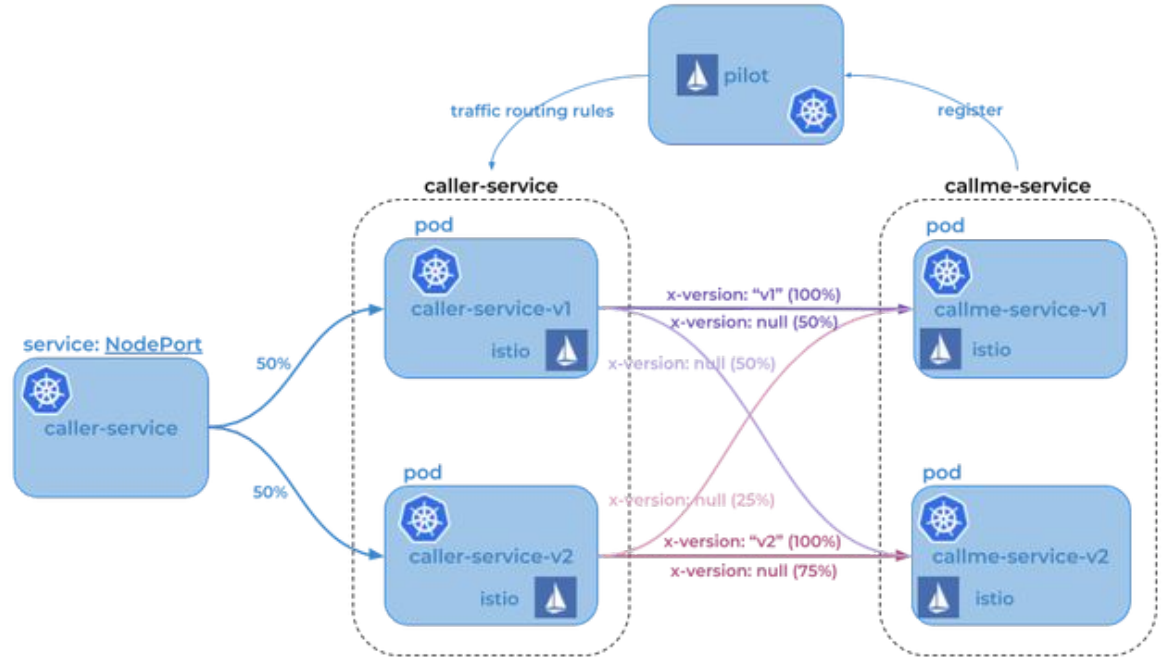
# Discovery and load balancing

- K8s DNS
- Istio virtual services

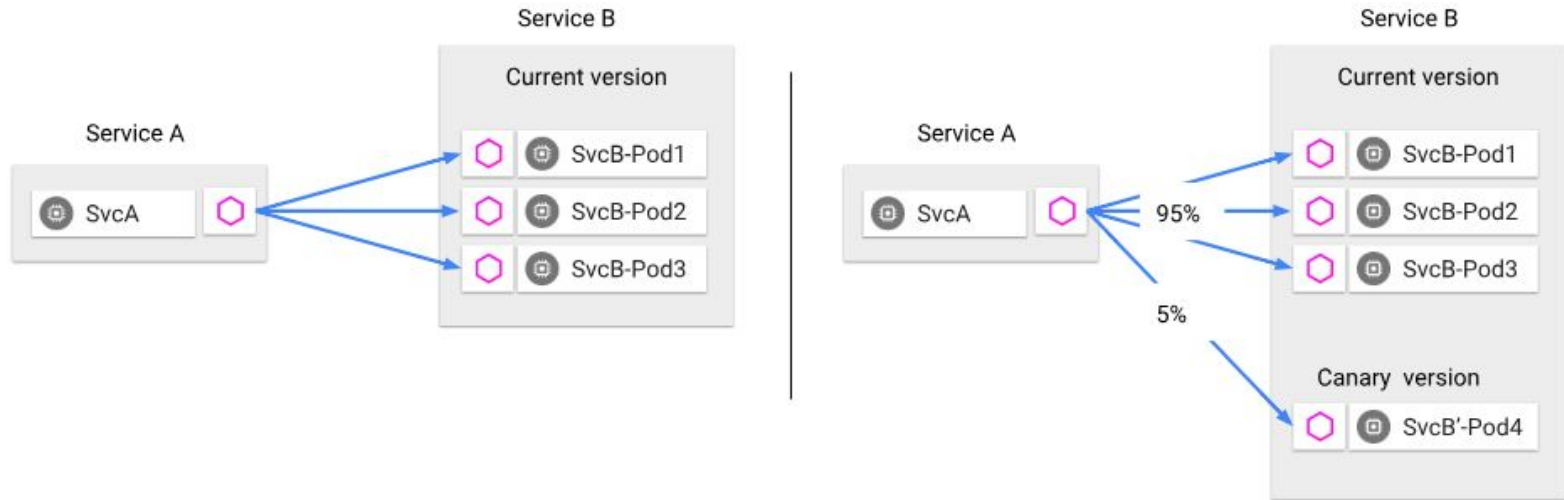


# Traffic management

- Routing rules
- Traffic shifting
- Traffic mirroring
- Conditional rules
- Fault ingestion

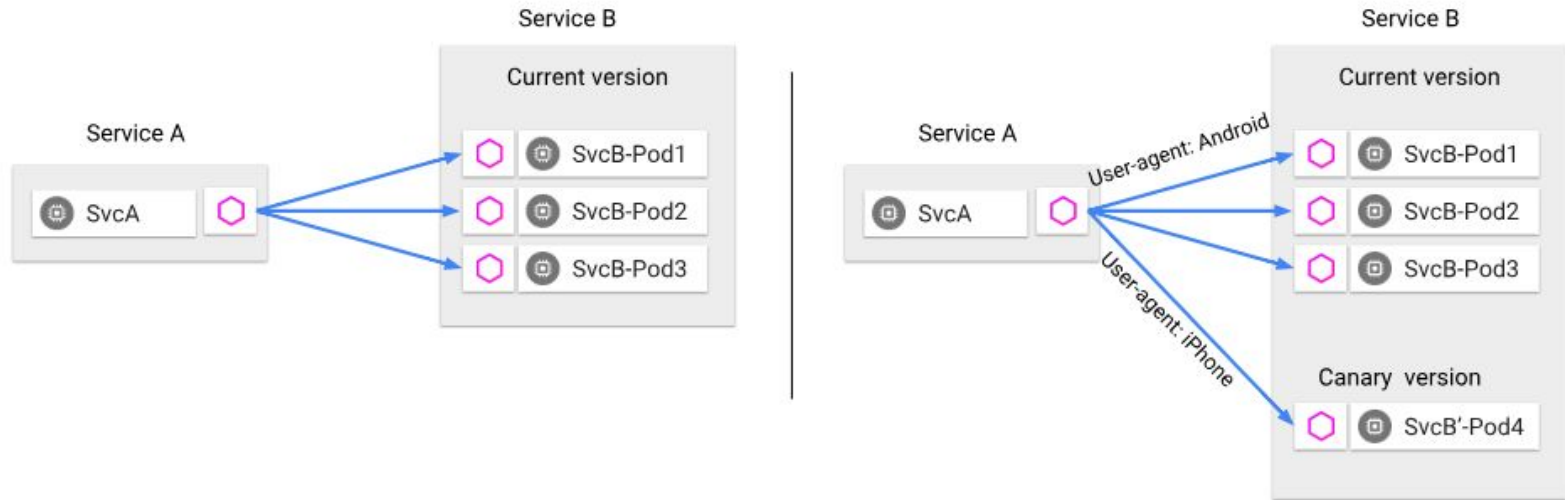


# Traffic management



**Traffic splitting decoupled from infrastructure scaling** - proportion of traffic routed to a version is independent of number of instances supporting the version

# Traffic management



**Content-based traffic steering** - The content of a request can be used to determine the destination of a request



# Handling failures

Envoy provides a set of out-of-the-box opt-in failure recovery features that can be taken advantage of by the services in an application.

Features include:

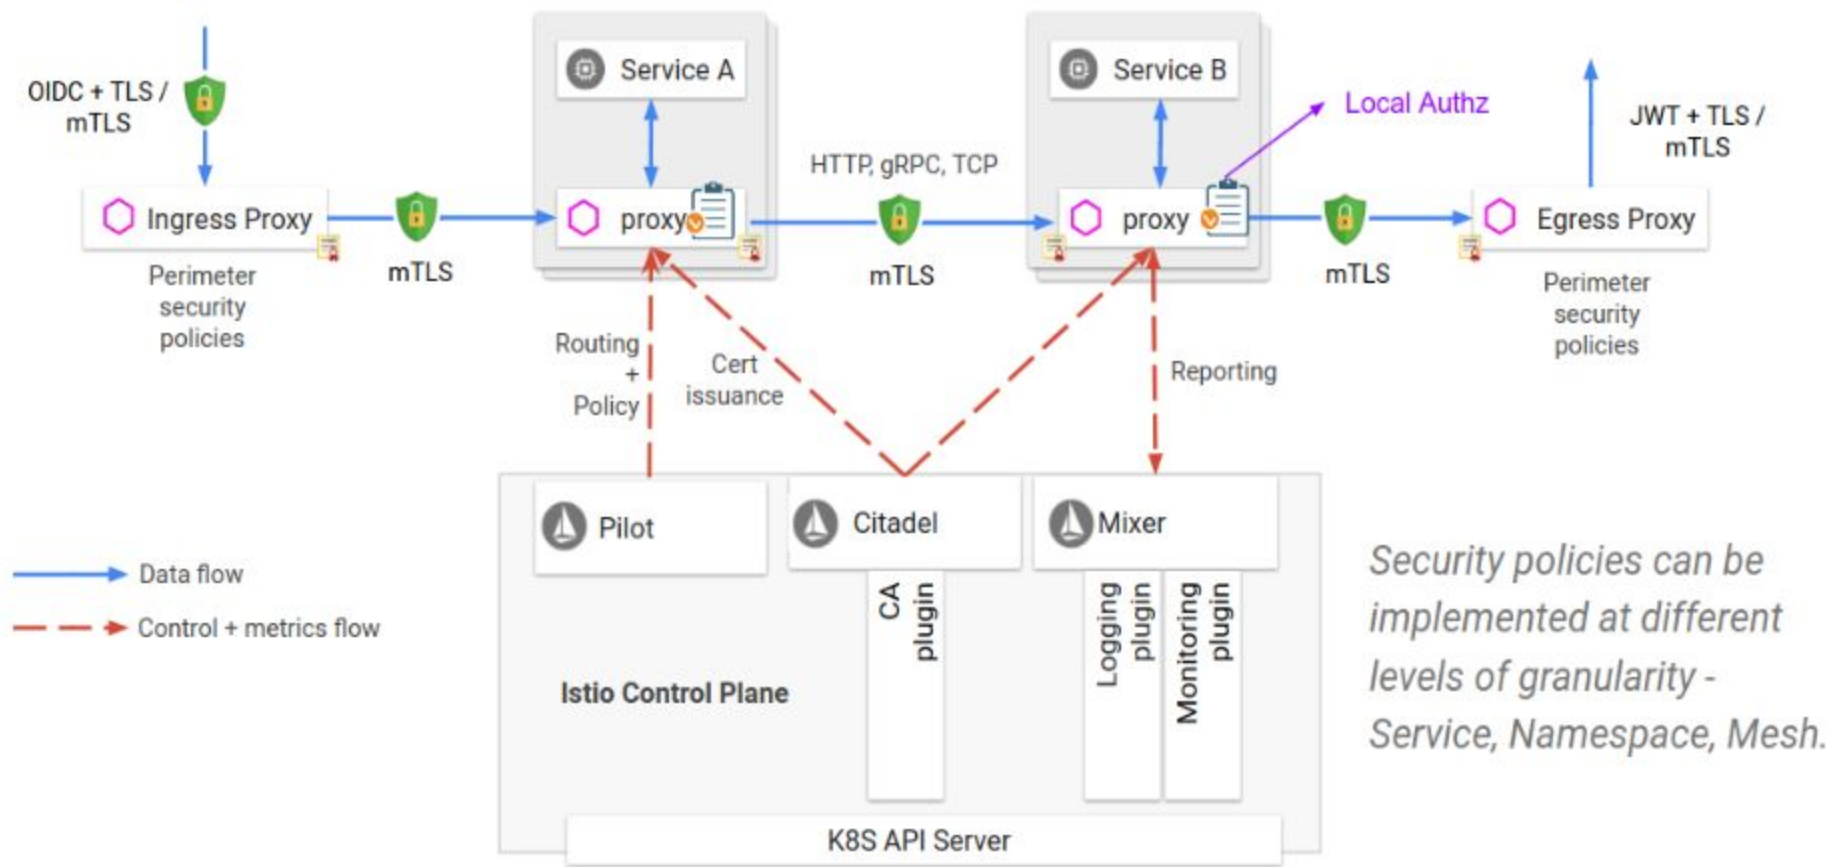
- Timeouts
- Bounded retries with timeout budgets and variable jitter between retries
- Limits on number of concurrent connections and requests to upstream services
- Active (periodic) health checks on each member of the load balancing pool
- Fine-grained circuit breakers (passive health checks) – applied per instance in the load balancing pool



# Istio Security

- Secure by default
- Zero-trust network
- Defense in depth
  
- Identity
- Policy
- Encryption
  
- Communication
- Endpoints
- Data
- Platform

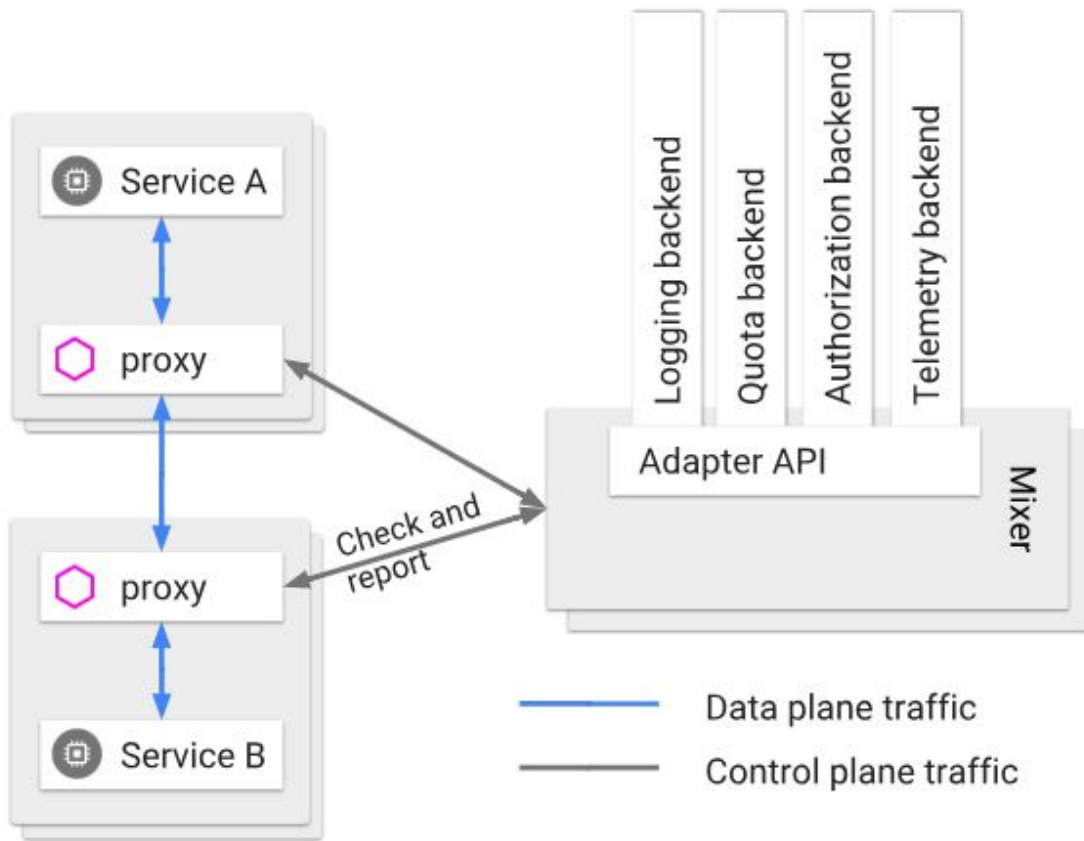




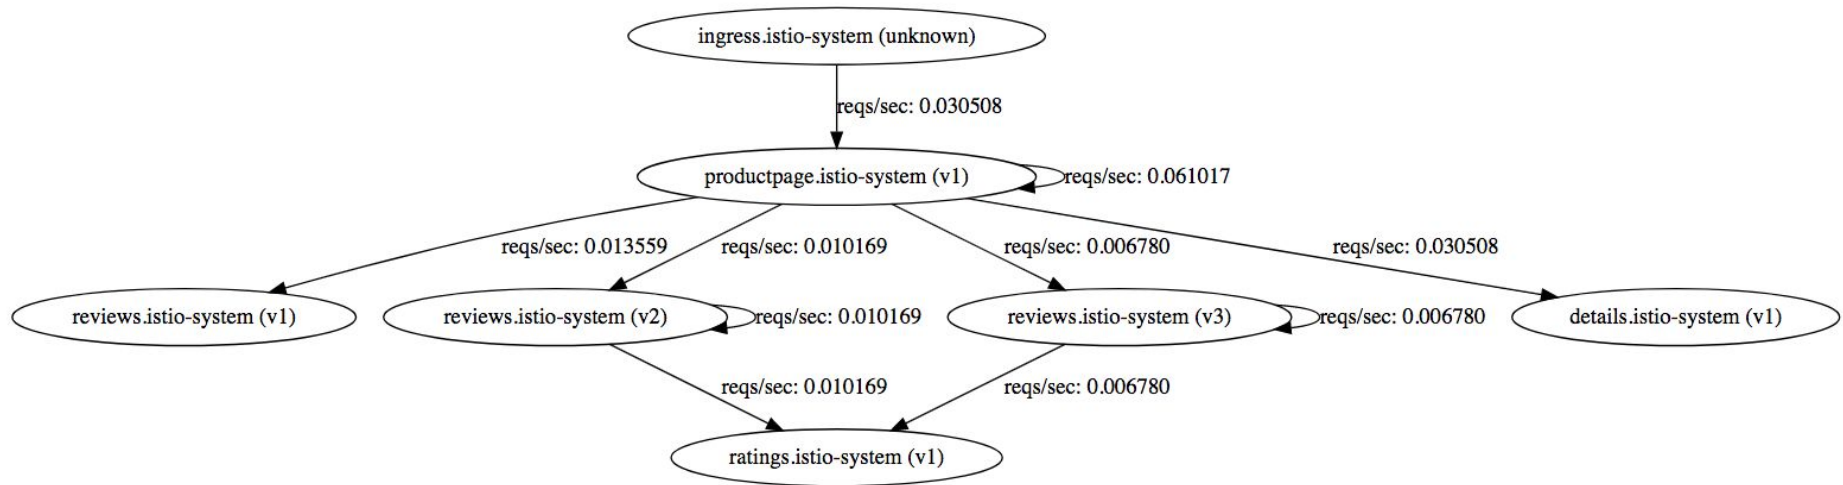
*Istio Security Architecture*

# Istio Observability

- Telemetry
- Logging
- Auditing



# Istio service graph



## Tracing with Jaeger

## Find Traces

Service

istio-ingress ▾

all ▾

Tags ⓘ

http.status\_code=200 error=true

Lookback

Last Hour ▴ ▾

Min Duration

e.g. 1.2s, 100ms, 5

Max Duration

e.g. 1.1s

Limit Results

200

Find Traces



200 Traces

Sort

Most Recent ▴ ▾

istio-ingress: productpage-default

37.09ms

8 spans

details (1) istio-ingress (1) productpage (3) ratings (1) reviews (2)

02:21:58 pm (a few seconds ago)

istio-ingress: productpage-default

39.43ms

8 spans

details (1) istio-ingress (1) productpage (3) ratings (1) reviews (2)

02:21:43 pm (a few seconds ago)

istio-ingress: productpage-default

40.08ms

8 spans

details (1) istio-ingress (1) productpage (3) ratings (1) reviews (2)

02:21:27 pm (a minute ago)

Max Duration

istio\_request\_count

Load time: 794ms  
Resolution: 1s  
Total time series: 10

Execute

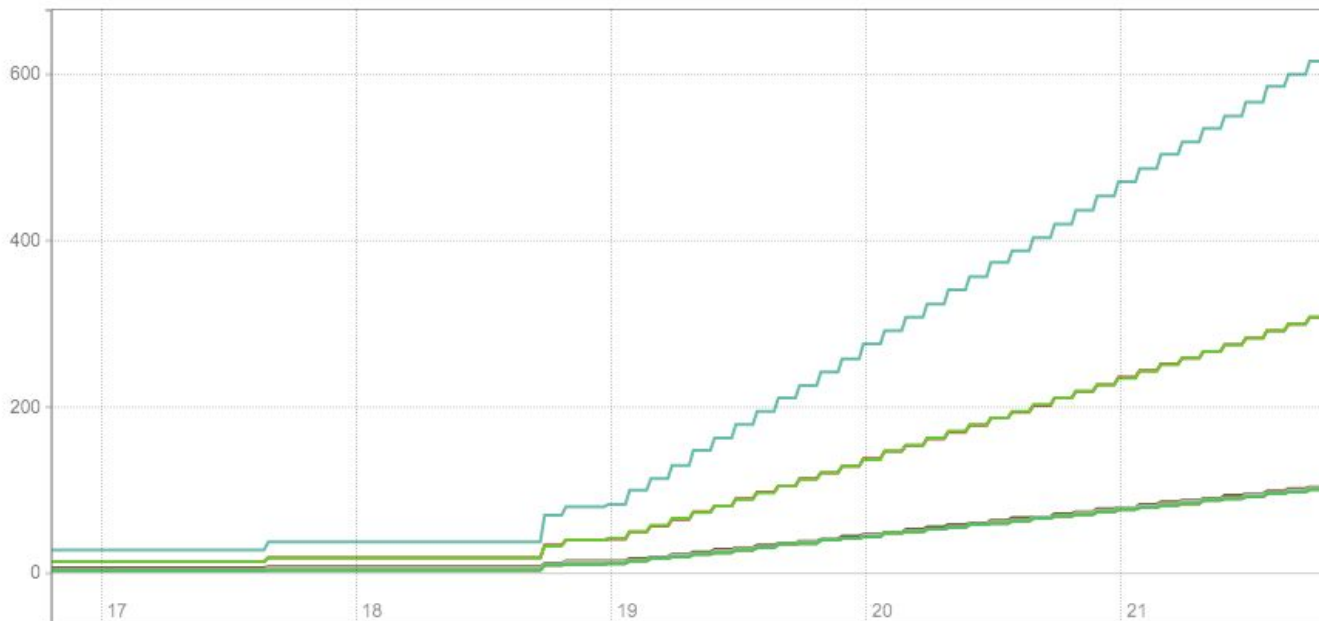
- insert metric at cursor -

istio\_request

Graph

Console

5m    ⏪    Until    ⏩    Res. (s)    ☐ stacked



Load time: 149ms  
Resolution: 14s  
Total time series: 10

Value

154

154

308

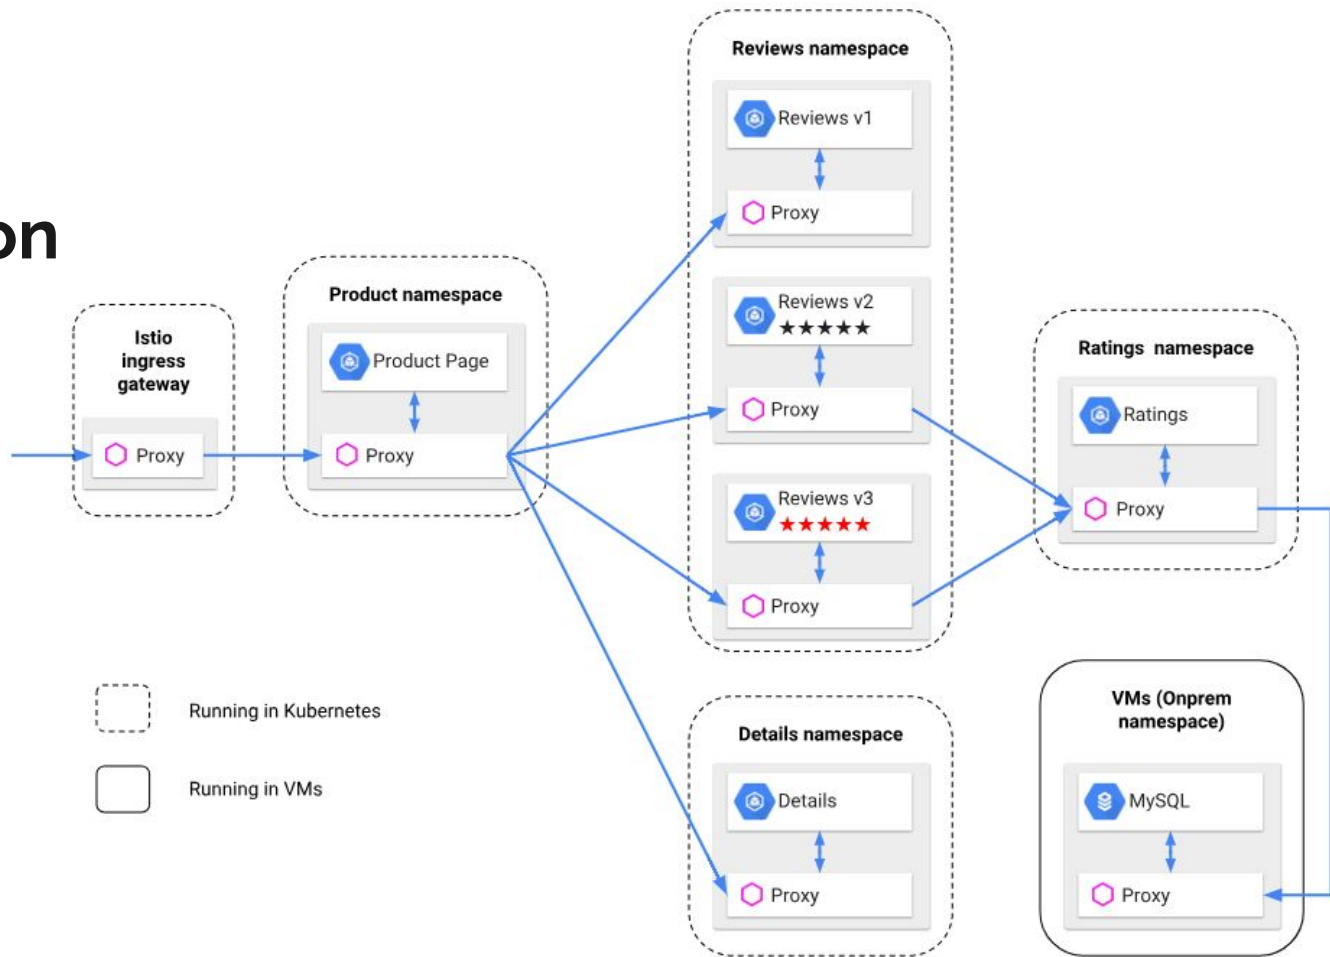
51

50

istio\_request\_count{destination\_service="reviews.istio-system.svc.cluster.local",destination\_version="v3",instance="istio-mixer.istio-system:42422",job="istio-mixer"}  
istio\_request\_count{destination\_service="reviews.istio-system.svc.cluster.local",destination\_version="v3",instance="istio-mixer.istio-system:42422",job="istio-mixer"}  
istio\_request\_count{destination\_service="reviews.istio-system.svc.cluster.local",destination\_version="v2",instance="istio-mixer.istio-system:42422",job="istio-mixer"}

# VM expansion

- Exposing Pilot, Mixer and Citadel outside of k8s cluster
- Envoy on VM
- Iptables rules





## Recap

- Open Source
- Production-Grade Service mesh
- Scales your control, not your ops team
- Runs anywhere (on-prem, any public cloud)
- Istio Features
  - Service discovery and load balancing
  - Black belt routing capabilities
  - Failure handling
  - Failure ingestion



# Istio

Connect, secure, control, and observe services.

- Service and user Identities + RBAC
- mTLS
- Tracing
- Telemetry
- Central logging



# Links

[Istio.io](https://istio.io) (official site)

[Microservices in the Cloud with Kubernetes and Istio \(Google I/O '18\)](#) (presentation)

[Lightning-Talk Style Demo of Istio and OpenCensus](#) (demo project on github)





**That's all folks!**  
**Q&A time**

Feel free to reach out at  
[linkedin.com/in/plstoyanov](https://www.linkedin.com/in/plstoyanov)