

# Kubernetes Security and Hardening

Kubernetes/CNCF Sofia Meetup  
January 2019

Orlin Vasilev



kubernetes

# \$ whoami - Orlix



Orlin Vasilev

- Software Infrastructure Engineer at VMware
- father of two
- snowboarder/DIY person/home BBQ pitmaster

twitter: @OrlinVasilev

GitHub: @OrlinVasilev

# Agenda



- Container Security
- Kubernetes Security
- Demo – kube-bench
- Q&A

# Goals of this talk



- **Raise** awareness of high-risk attacks possible in many installs
- Demonstrate few attacks
- Provide hardening methods
- Share additional hardening tips

# Possible Attack Surfaces



**Underlying Infrastructure** - OS , Metadata API

Kube-apiserver

**Container** - Malicious Image, Compromised Container

Dashboard Access

**Escape the container**

Direct Etcd Access

**Metrics Scraping** - read-only ports

Application Tampering

**Kubelet**

Launch too many pods / high consume CPU/RAM/Disk

**Service Account Tokens**

Docker daemon security(or containerd, rkt, CRI-O ...)



kubernetes

# Container Security



Container Runtime - Least Privileges

Base Image

Image Builder/Maintainer

Image Scanning

Image Signing

**TIP: do not run as root in container**

```
RUN useradd -r -u 1001 -g appuser appuser  
USER appuser
```



kubernetes

# Few Kubernetes Security Aspects



- Kubernetes security:
  - Properly configured RBACs
  - Secrets
  - Pod Security Policy
  - Network Policy
  - Admission Controllers
  - etc...

# Kubernetes RBAC



- Enabled in kube-api : `—authorization-mode=RBACRole` and `ClusterRole`
- `RoleBinding` and `ClusterRoleBinding`
- `ServiceAccounts` and `User Accounts`

## Tip:

use `kube-apiserver` audit logs to monitor activities



# Role and RoleBinding



```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: ["" ] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: jane # Name is case sensitive
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role #this must be Role or ClusterRole
  name: pod-reader #must match name
  apiGroup: rbac.authorization.k8s.io
```



# Kubernetes Secrets



- A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key.
- base64ed
- consumed within a pod as volumeMounts
- secrets used to keep docker login information for private registry
- use secrets instead of writing sensitive data in the containers

```
apiVersion: v1
data:
  .dockerconfigjson: eyJhdXRocyI6eyJk...(long base64)
kind: Secret
metadata:
  name: regcred
  namespace: secconf
type: kubernetes.io/dockerconfigjson
```



# PodSecurityPolicy



- Enforced by the PodSecurityPolicy admission controller enabled on kube-api
- PSP is an ClusterLevel resource
- Set of conditions which allows a pod to be run
- some examples
  - [privileged](#)
  - [runAsUser](#)
  - [hostPID](#)
  - [volumes](#)

## Tips:

- create PSP before enabling PSP admission controller
- PSP are applied alphabetically

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: privileged
  annotations:
    seccomp.security.alpha.kubernetes.io/allowedProfileNames: '*'
spec:
  privileged: true
  allowPrivilegeEscalation: true
  allowedCapabilities:
  - '*'
  volumes:
  - '*'
  hostNetwork: true
  hostPorts:
  - min: 0
    max: 65535
  hostIPC: true
  hostPID: true
  runAsUser:
    rule: 'RunAsAny'
  selinux:
    rule: 'RunAsAny'
  supplementalGroups:
    rule: 'RunAsAny'
  fsGroup:
    rule: 'RunAsAny'
```

Usage via RBAC with ClusterRoles and ClusterRoleBindings to ServiceAccount or User



kubernetes

# Network Policy



A network policy is a specification of how groups of pods are allowed to communicate with each other and other network endpoints.

Uses labels to select pods and define rules.

NetworkPolicy needs network plugin which supports it.(calico/weave..)

By default, if no policies exist in a namespace, then all ingress and egress traffic is allowed to and from pods in that namespace.

# NetworkPolicy example



```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
            except:
              - 172.17.1.0/24
```

```
- namespaceSelector:
    matchLabels:
      project: myproject
- podSelector:
    matchLabels:
      role: frontend
ports:
  - protocol: TCP
    port: 6379
egress:
  - to:
      - ipBlock:
          cidr: 10.0.0.0/24
    ports:
      - protocol: TCP
        port: 5978
```



# CIS Kubernetes Benchmark



- 200+ pages of best practices and tests  
<https://www.cisecurity.org/benchmark/kubernetes/>
- Tests for workers, masters and federated nodes
- **kube-bench** (Aqua Security)
  - 1265 stars on GitHub
  - ~30 contributors
  - <https://github.com/aquasecurity/kube-bench>



# more tools



- **sonobuoy**(heptio)
  - 1141 stars on GitHub
  - 33 contributors
  - <https://github.com/heptio/sonobuoy>
- **kubeaudit** (shopify)
  - 306 stars on GitHub
  - 16 contributors
  - <https://github.com/Shopify/kubeaudit>
- **k8sguard**
  - 122 stars on GitHub
  - 8 contributors
  - <https://github.com/k8sguard/k8sguard-start-from-here>



# Demo

Demo 1 - scraping unsecured metrics

Demo 2 - misconfigured taints for master and exploiting the lack of PSP

Demo 3 - NetworkPolicy





aliases:

k = kubectl

kdp = kubectl describe pods

kep = kubectl edit pod

kgn = kubectl get nodes —show-lables -o wide

lp = kubectl get pods

lpa = kubectl get pods —all-namespaces

kshell = kubectl exec -it \$\$ bash

kn = kubens



### General Guidance

1. Verify that all your security settings properly enforce the policy
2. Use the latest stable K8s version possible to gain the latest security capabilities and fixes
3. Audit the OS, container runtime, and K8s configuration using CIS Benchmarking and other tools like [kube-auto-analyzer](#) and [kube-bench](#)
4. Log **everything** to a location **outside** the cluster

### Image Security

1. Use private registries, and restrict public registry usage
2. Scan all images for security vulnerabilities continuously. E.g [CoreOS Clair](#) or [Atomic Scan](#)
3. Decide which types/severity of issues should prevent deployments
4. Maintain standard base images and ensure that all workloads use them
5. **Do NOT run containers as the root user**

## K8s Components Security

1. API Server *authorization-mode=Node,RBAC*
2. Ensure all services are protected by TLS
3. Ensure *kubelet* protects its API via *authorization-mode=Webhook*
4. Ensure the *kube-dashboard* uses a restrictive *RBAC* role policy and v1.7+
5. Closely monitor all *RBAC* policy failures
6. Remove default *ServiceAccount* permissions

## Network Security

1. Filter access to the cloud provider metadata APIs/ URL, and Limit IAM permissions
2. Use a CNI network plugin that filters ingress/ egress pod network traffic
  - a. Properly label all pods
  - b. Isolate all workloads from each other
  - c. Prevent workloads from egressing to the Internet, the Pod IP space, the Node IP subnets, and/or other internal networks
  - d. Restrict all traffic coming into the kube-system namespace except kube-dns
3. Consider a Service Mesh!

## Hardening Tips (Continued)

### Workload Containment and Security

1. Namespaces per tenant
2. Default network “deny” inbound on all namespaces
3. Assign CPU/RAM *limits* to all containers
4. Set *automountServiceAccountToken: false* on pods where possible
5. Use a *PodSecurityPolicy* to enforce container restrictions and to protect the node
6. Implement container-aware malicious activity / behavioral detection

### Misc Security

1. Collect logs from all containers, especially the RBAC access/deny logs
2. Encrypt the contents of *etcd*, and run *etcd* on dedicated nodes
3. Separate Cloud accounts/VPCs/projects/resource groups
4. Separate clusters for dev/test and production environments
5. Separate node pools for different tenants

# Resources and goodreads



- <https://github.com/kelseyhightower/kubernetes-the-hard-way>
- <https://github.com/hardening-kubernetes/from-scratch>
- <https://github.com/cncf/presentations/tree/master/>
- <https://goo.gl/TNRxtd>





# Q&A