

DBaaS: An autonomous journey in a multi-cloud Kubernetes environment

SCAN ME



About me:

```
{  
  Title: "Staff Engineer",  
  Profile: "Golang",  
  Env: "Multi-Cloud",  
  Platform: "Kubernetes"  
}
```

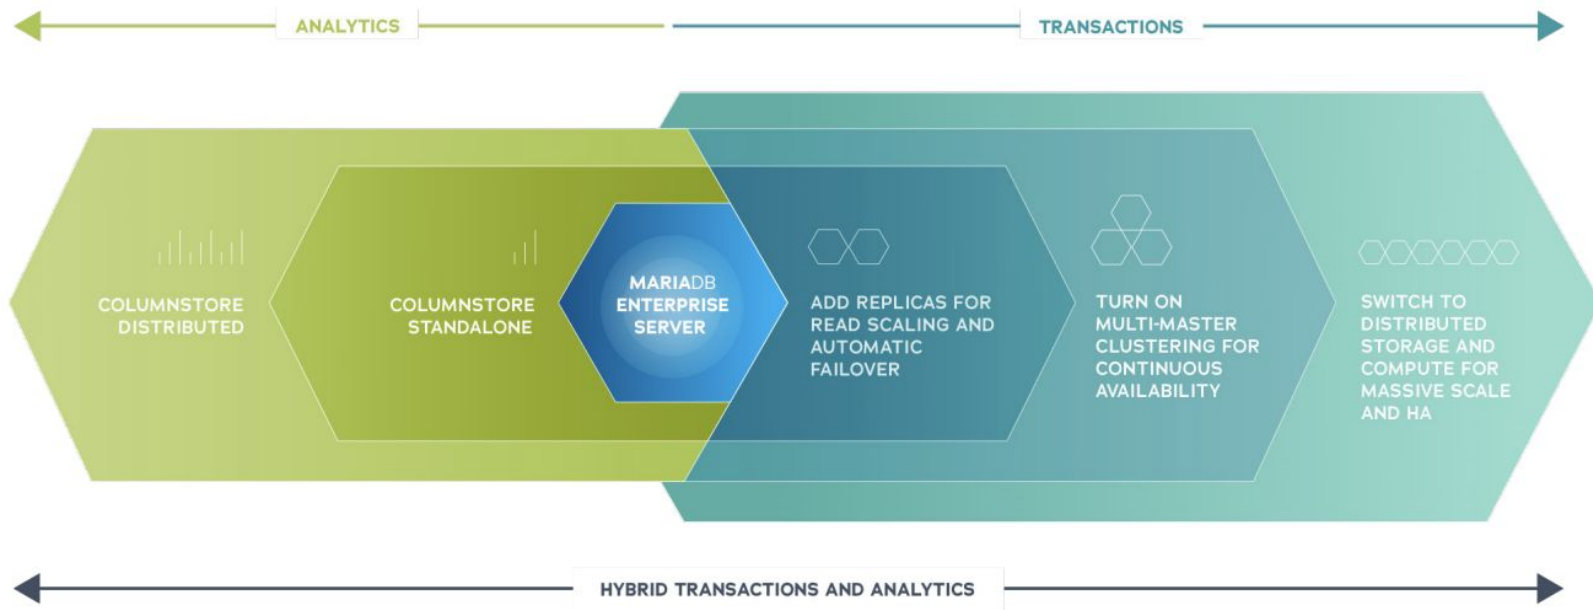
What is SkySQL

- fully-managed cloud database service
- deploys and manages databases for the customer with a few clicks
- Platform offers:
 - Enterprise Server: Single and Primary/Replica
 - Xpand - Distributed SQL: Multi-primary distributed SQL
 - Column Store
 - Serverless Analytics
 - Observability Services
 - Multicloud deployment GCP/AWS

SCAN ME



SkySQL Products Overview



Customer Showcase: SAMSUNG

More than 2 Billion smart phones sold.

All SAMSUNG Android phones rely on Xpand to access the SAMSUNG Cloud.

Autonomous DBaaS

Create - Use

SkySQL will take care of it...

Autonomous DBaaS

- Disk Autoscaling
 - Based on current disk utilization
 - Prediction based on Prometheus [predict_linear\(\)](#) over some interval (Linear Regression)
- Horizontal Scaling
 - CPU utilisation over all replicas sustained for some interval
 - Number of concurrent sessions over all replicas sustained for some interval
 - Prediction: Number of concurrent sessions is expected to hit the boundary within 4h over on the last hour data
- Vertical Scaling
 - CPU & Memory utilisation over all replicas sustained for some interval
 - Prediction on CPU & Memory utilization

Autonomous DBaaS

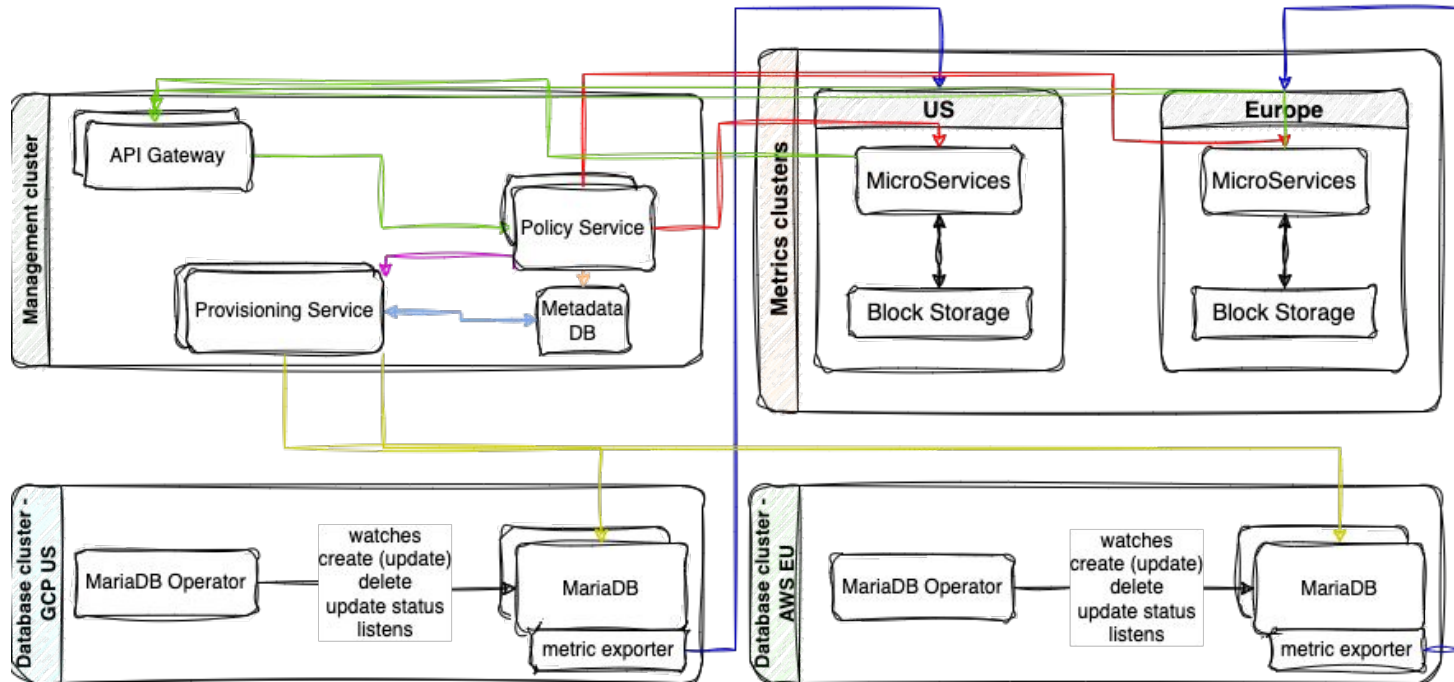
- SkySQL offers Terraform provider
- allows configuring any SkySQL DB topology using the Terraform's declarative language
- <https://registry.terraform.io/providers/mariadb-corporation/skysql/latest/docs>

```
# Create a service
resource "skysql_service" "default" {
  service_type = "transactional"
  topology     = "es-single"
  cloud_provider = "gcp"
  region       = "us-central1"
  name         = "myservice"
  architecture = "amd64"
  nodes        = 1
  size         = "sky-2x8"
  storage      = 100
  ssl_enabled  = true
  version      = data.skysql_versions.default.versions[0].name
  # [Optional] Below you can find example with optional parameters how to configure a privatelink connection
  endpoint_mechanism = "privatelink"
  endpoint_allowed_accounts = ["gcp-project-id"]
  # [/Optional]
  # The service create is an asynchronous operation.
  # if you want to wait for the service to be created set wait_for_creation to true
  wait_for_creation = true
  # You need to add your ip address in the CIRD format to allow list in order to connect to the service
  allow_list = [
    {
      "ip" : "104.28.203.45/32",
      "comment" : "office"
    }
  ]
}
```


MariaDB Operator

- SkySQL offers enterprise kubernetes operator
- Kubebuilder <https://github.com/kubernetes-sigs/kubebuilder>
- Developed on top of `controller-runtime` and `controller-tools` libraries
- *Controller-gen* generates CRDs, RBACs manifests etc, as well generate go lang deepcopy methods for apimachinery interface
- Set of Controllers, which handles
 - Disk resizing, IOPs etc
 - Horizontal/Vertical scaling
 - Backups/Restore
 - Configuration Management

Autonomous Flow: multi cloud/region



Not the kubectl way

- Get Cluster via Cloud SDK
 - CA
 - Cluster endpoint
 - Access Token
- Connect, take care of NAT,VPC

```
caCertificate, err := base64.StdEncoding.DecodeString(cluster.CaCertificate)
if err != nil {
    return errors.Wrap(err, "not able to decode cluster ca certificate")
}

tlsClientConfig := rest.TLSClientConfig{}
tlsClientConfig.CAData = caCertificate
config := &rest.Config{
    Host:      cluster.Endpoint,
    TLSClientConfig: tlsClientConfig,
    BearerToken: cluster.AccessToken,
}

scheme := runtime.NewScheme()
if err = mdbv1alpha1.AddToScheme(scheme); err != nil {
    return errors.Wrapf(err, "could not add mariadb kubernetes to schema")
}
if err = corev1.AddToScheme(scheme); err != nil {
    return errors.Wrapf(err, "could not add corev1 kubernetes to schema")
}

// Create the client
clientSet, err := client.New(config, client.Options{
    Scheme: scheme,
})
if err != nil {
    return err
}


mdbList := &mdbv1alpha1.MariaDBList{}
if err = clientSet.List(ctx, mdbList); err != nil {
    return err
}

log.Println("Available MariaDBs for this cluster:")
for _, ns := range mdbList.Items {
    log.Printf("*\t Mariadb Name %s, Namespace %s", ns.Name, ns.Namespace)
}
```

Multi-Cloud Challenges

- GCP Golang sdk is structured significantly better than AWS
- GCP doesn't allow update of existing node pool, creates new one and cordones
- AWS supports updating node pool
- GCP latency between the zones ridiculously lower, similar to single zone latency in AWS

DEMO: Disk Autoscaling



SCAN ME



Ask me Anything
