

# Revolutionizing Cloud Infrastructure: Unikraft and Unikraft Cloud - Unleashing the Power of Unikernels

---

Todor Todorov

Staff Engineer, Man Group UK Ltd.

# Agenda

- What is a Unikernel?
- Introduction to Unikraft
- Introduction to Unikraft Cloud
- Demo
- Q&A

# Agenda

- What is a Unikernel?

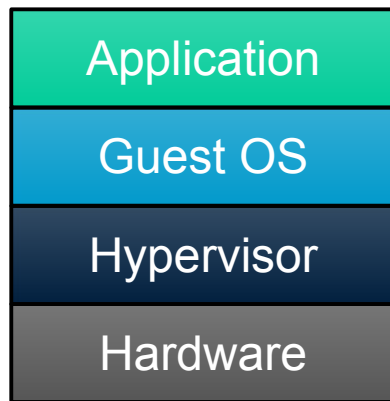
# What is a Unikernel?

- Specialized library operating system
- Single process\*
- Multi-threading
- Single user
- Single address space


# Benefits

- Reduced resource consumption
- Enhanced security
- Faster boot times
- Improved performance
- Simplified deployment

# Traditional VM vs Unikernel



Traditional  
Virtual Machine



Unikernel

# Notable Projects

- <http://unikernel.org/projects/>
- OSv <https://osv.io/> (2012-12-01)
- MirageOS <https://mirage.io/> (2013-02-02)
- IncludeOS <https://www.includeos.org/> (2014-07-07)
- Hermit OS <http://hermit-os.org> (2015-05-09)
- Unikraft <https://unikraft.org/> (2017-11-29)

# Agenda

- Introduction to Unikraft



# Introduction to Unikraft

- Single address space
- Fully modular system
- Single protection level
- Static linking
- POSIX support
- Platform abstraction

# KraftKit

- Building, packaging, running locally and deploying unikernels to Unikraft Cloud
- Resembles the docker command
- Supports compose files
- Manages resources on Unikraft Cloud
- Can read instructions from Kraftfile
- Can set KConfig options

# Kraftfile

spec: v0.6

name: nginx

rootfs: ./Dockerfile

cmd: ["/usr/bin/nginx"]

template:

source: <https://github.com/unikraft/app-elfloader.git>

version: staging

unikraft:

source: <https://github.com/unikraft/unikraft.git>

version: staging

kconfig:

CONFIG\_LIBPOSIX\_ENVIRON\_ENV0: "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"

CONFIG\_LIBPOSIX\_ENVIRON\_ENV1: "LD\_LIBRARY\_PATH=/usr/local/lib:/usr/lib:/lib"

CONFIG\_LIBPOSIX\_ENVIRON\_ENV2: "HOME="/

CONFIG\_LIBPOSIX\_ENVIRON: 'y'

libraries:

lwip:

source: <https://github.com/unikraft/lib-lwip.git>

version: staging

kconfig:

CONFIG\_LWIP\_LOOPIF: 'y'

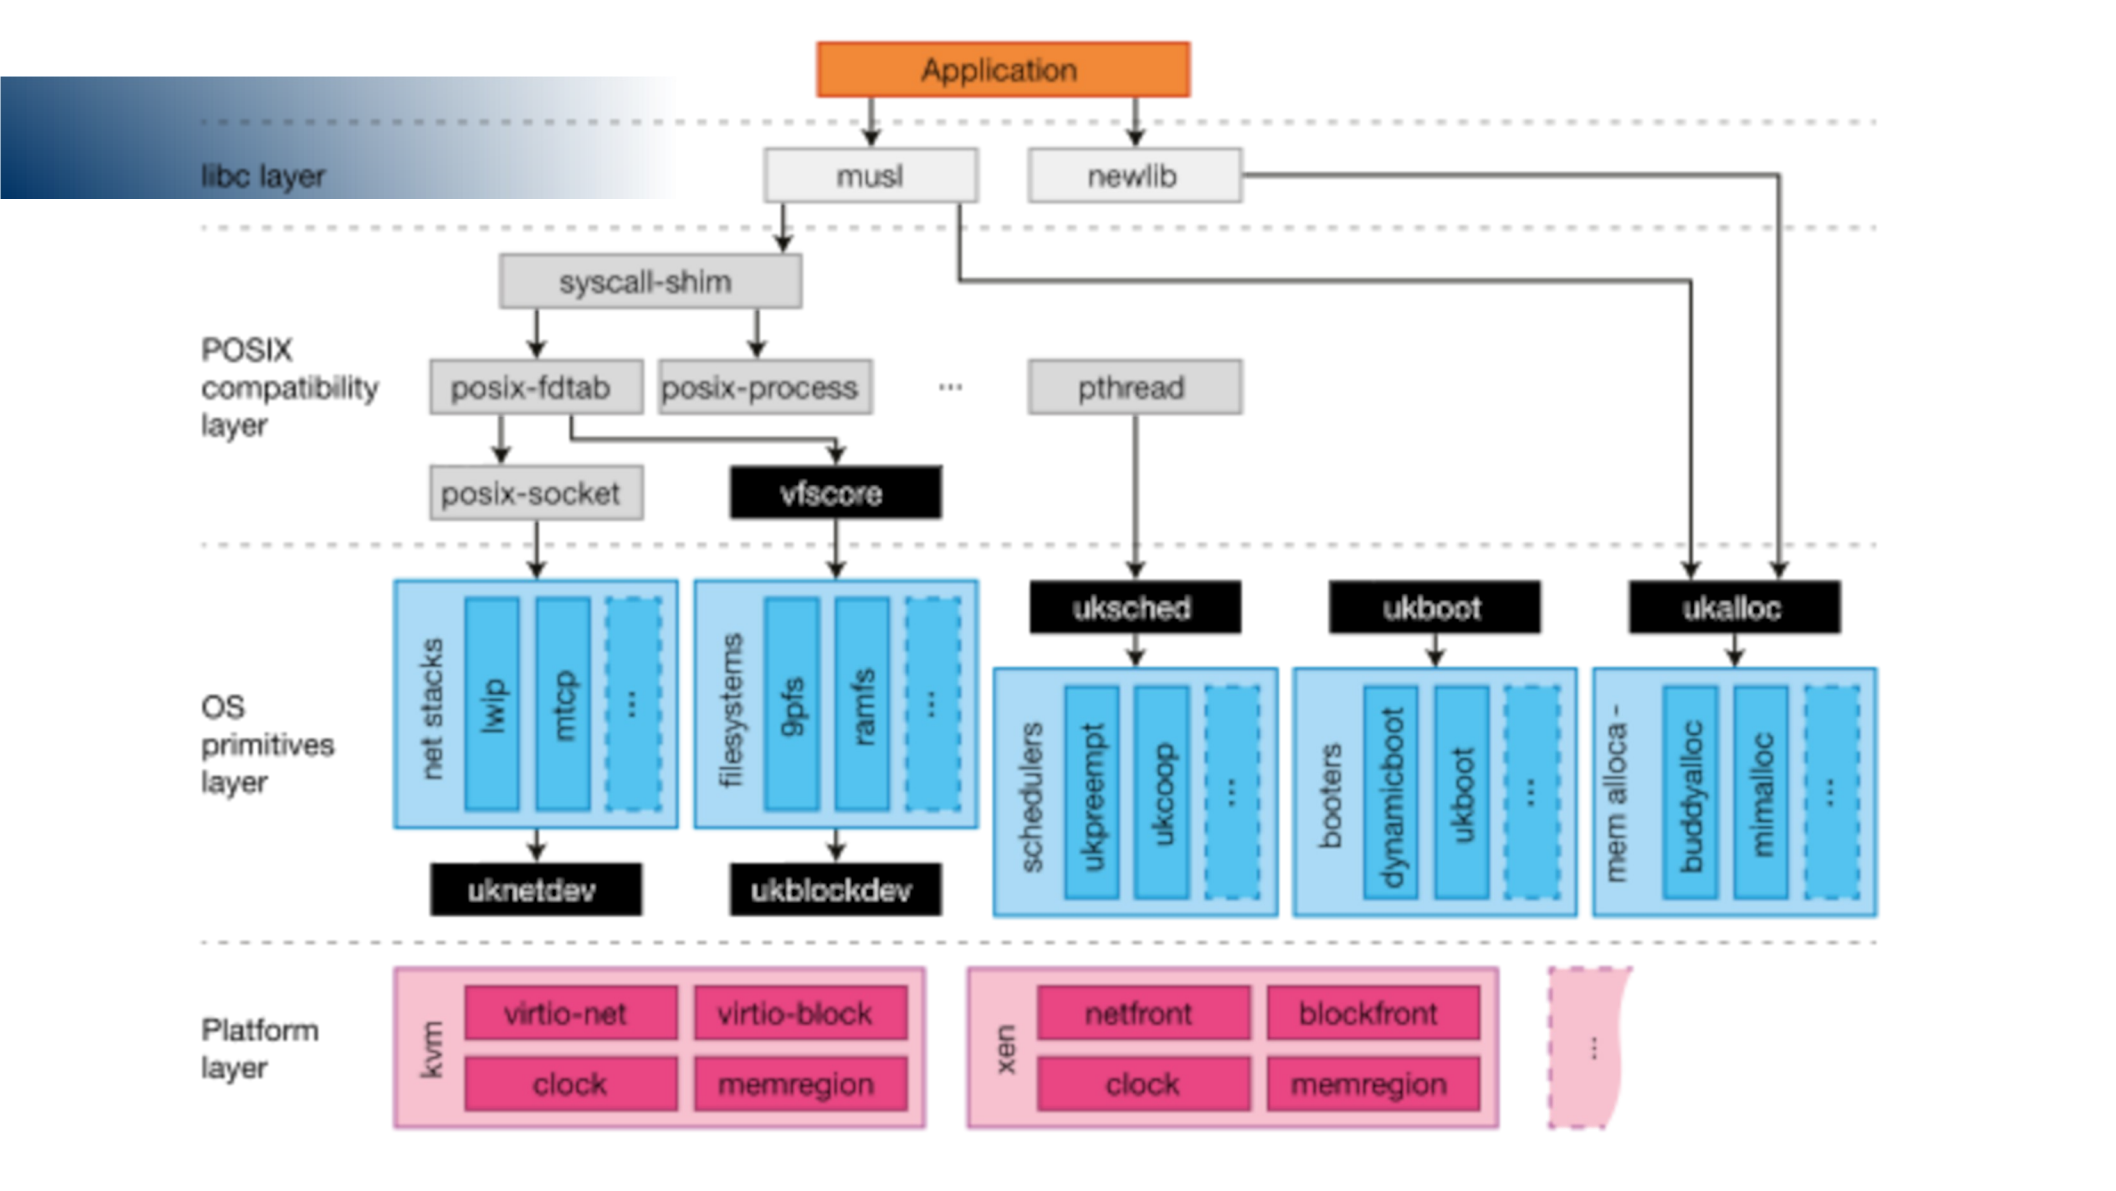
CONFIG\_LWIP\_UKNETDEV: 'y'

targets:

- fc/x86\_64
- qemu/x86\_64

# Architecture

- Micro-libraries – implement core Unikraft APIs
- Build system – Kconfig-based menu



# Build System

- Makefile – specifies locations of the Unikraft repository and external libraries; not needed for porting external libraries
- Makefile.uk – specifies sources, include paths, flags, and application-specific targets
- Config.uk – populates Unikraft's menu with application-specific options
- exportsyms.uk – lists symbols to be exported
- extra.ld – optional amendment to the main linker script

# Build Process

1. Configuring the Unikraft unikernel
2. Fetching and preparing the source code
3. Compiling the libraries and the core Unikraft code
4. Linking the final unikernel image

# Debugging

- Strace-like messages for binary system calls
- Global or per compilation unit debug messages
- Connecting GDB to QEMU's GDB server



# Agenda

- Introduction to Unikraft Cloud

# Introduction to Unikraft Cloud

- Based on unikernels
- Milliseconds boot time
- Auto-scale and scale to zero
- No traffic metering
- Based on Equinix
- Leveraging Firecracker as a Virtual Machine Monitor

# Concepts and Features

- Instance – running virtual machine
- Service – exposing the apps to the Internet
- All instances have a private IP and DNS by default
- Restart policy – never, always or on-failure
- Instances are created from images pulled from an OCI compatible registry
- Support of volumes, custom domains and rolling updates

# Integrations


- RESTful API
- GitHub Actions for building and deploying services
- Crossplane provider for deploying and managing unikernel instances as Kubernetes resources
- Terraform provider for deploying and managing unikernel instances
- metrics API exporting metrics in JSON or Prometheus format

# Millisecond Scale to Zero


User




# Millisecond Scale to Zero



# Millisecond Scale to Zero




# Millisecond Scale to Zero







# Millisecond Scale to Zero




# Millisecond Scale to Zero



# Millisecond Scale to Zero




# Millisecond Scale to Zero

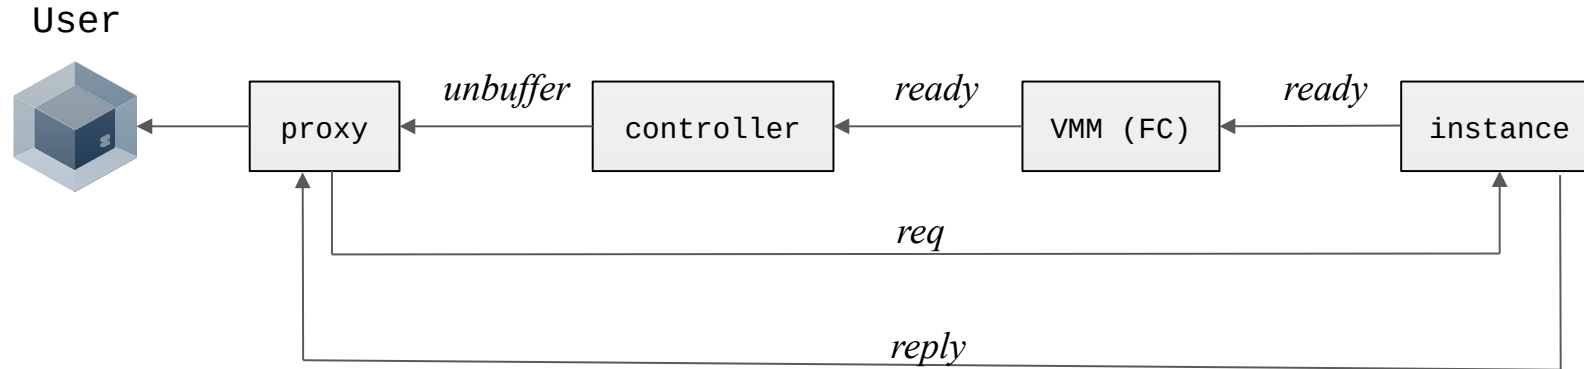


# Millisecond Scale to Zero

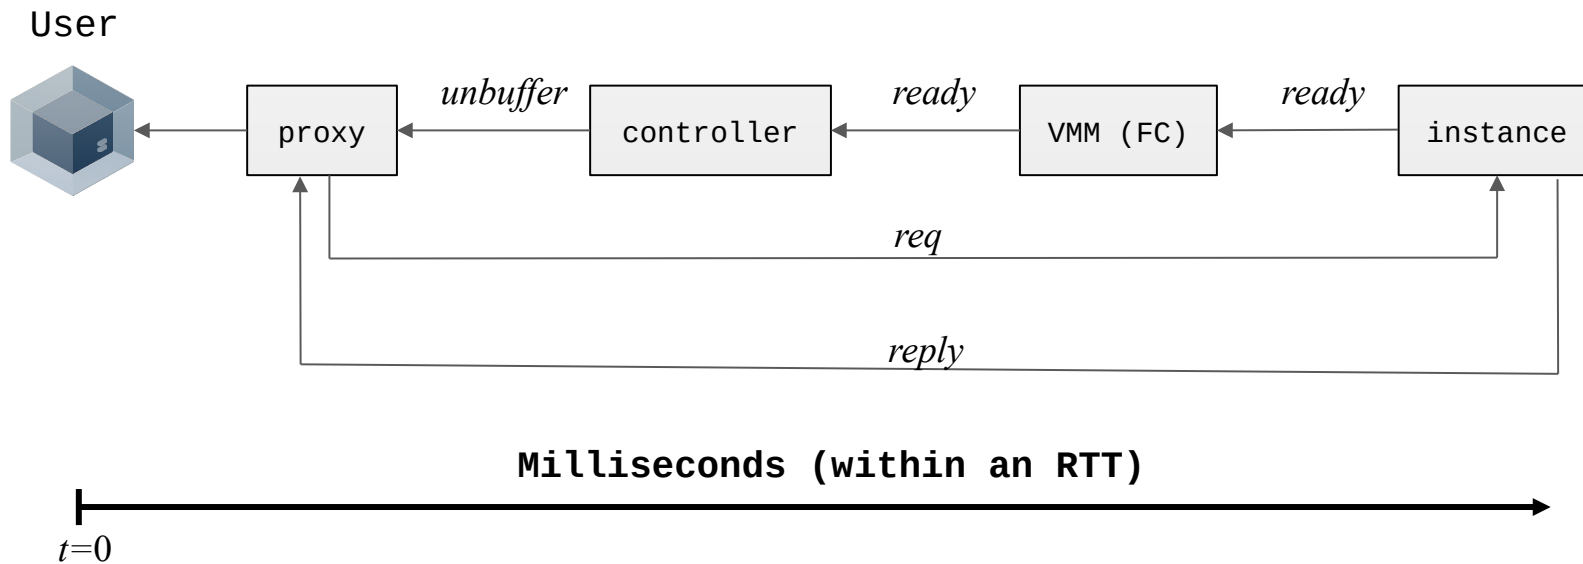
User



# Millisecond Scale to Zero




# Millisecond Scale to Zero



# Why it is Fast


- Network: custom load balancer = millis response time






# Why it is Fast

- Network: custom load balancer = millis response time
- Cloud stack: customer controller and images = millis response times



# Why it is Fast

- Network: custom load balancer = millis response time
- Cloud stack: customer controller and images = millis response times
- Application start time: snapshotting = millis response time



**DEMO**

**Q & A**

**Thank you**