

# Real-Time Data Processing with Apache Kafka: A Kubernetes-Native Approach






# #whoami






Iliyan Petkov

Valentin Hristev



 iliyanspetkov  
 @Iliyan\_petkov  
 @Iliyan-s-petkov



 valentin-hristev  
 @vhrstev

<https://dojobits.io/>

# What we do



**Professional  
services**



**Management  
services**



**Cloud design  
and  
architecture**




**Trainings**



# Our Goals:

Define the  
Problem



Apache  
Kafka  
Overview




Running Apache  
Kafka on  
Kubernetes

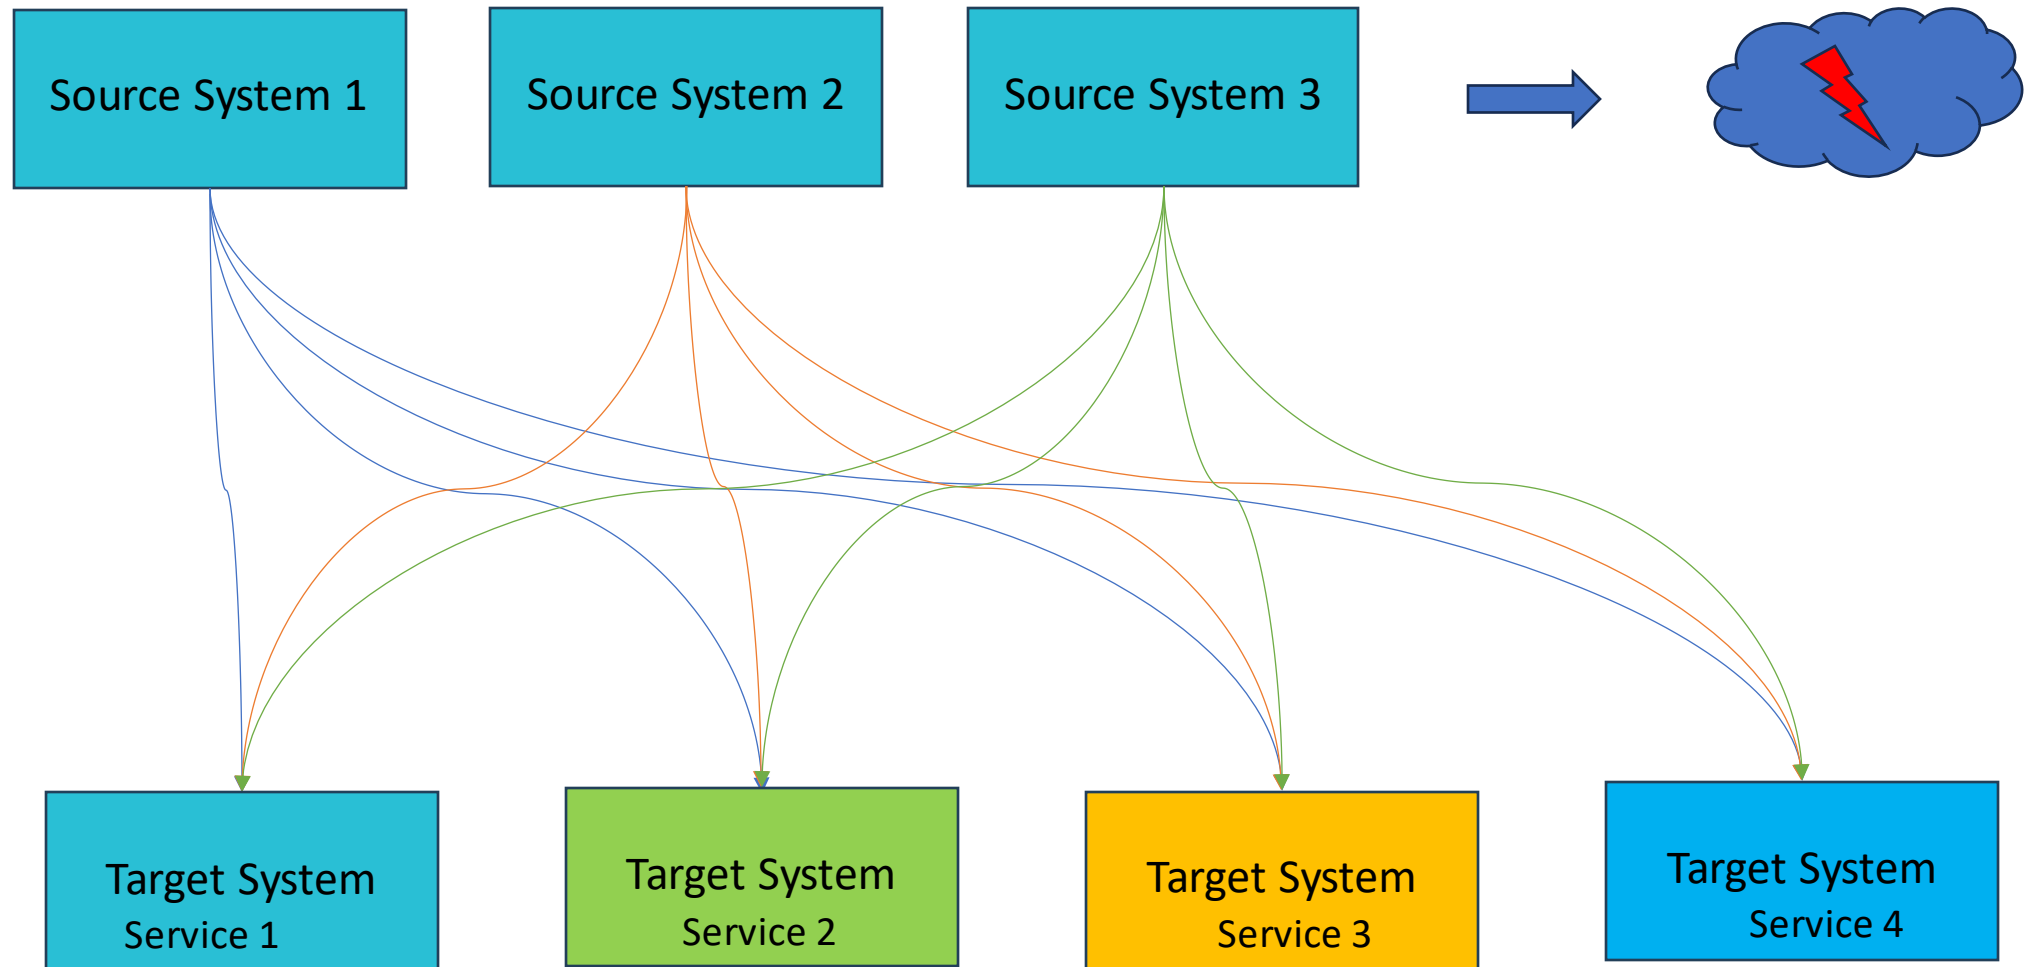


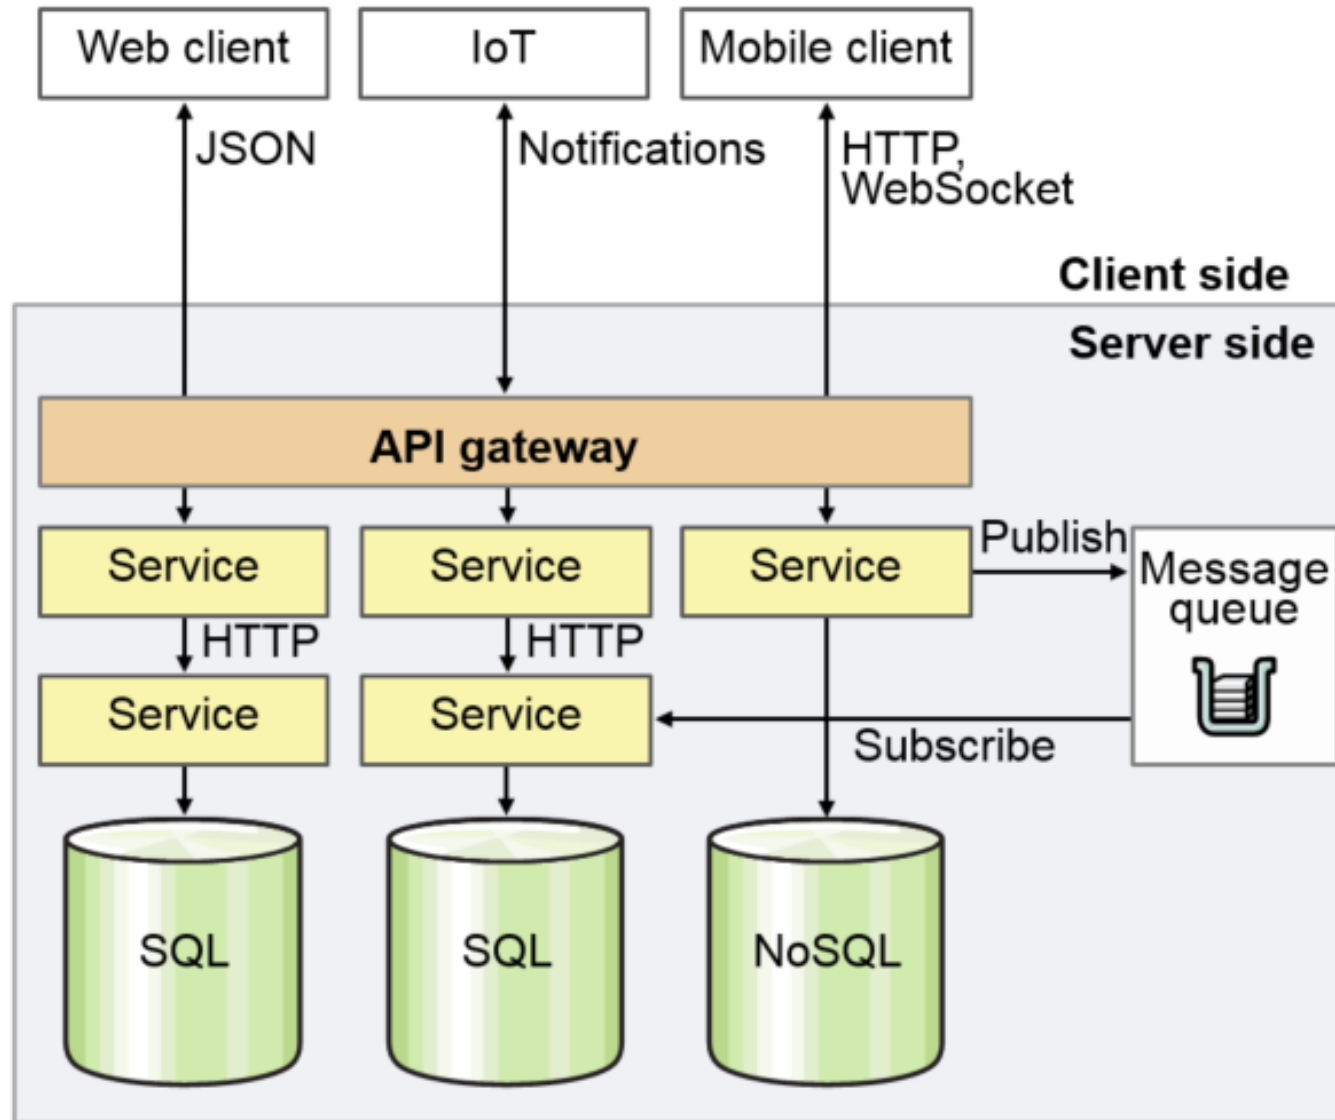
Never show a  
single line of  
JAVA code...

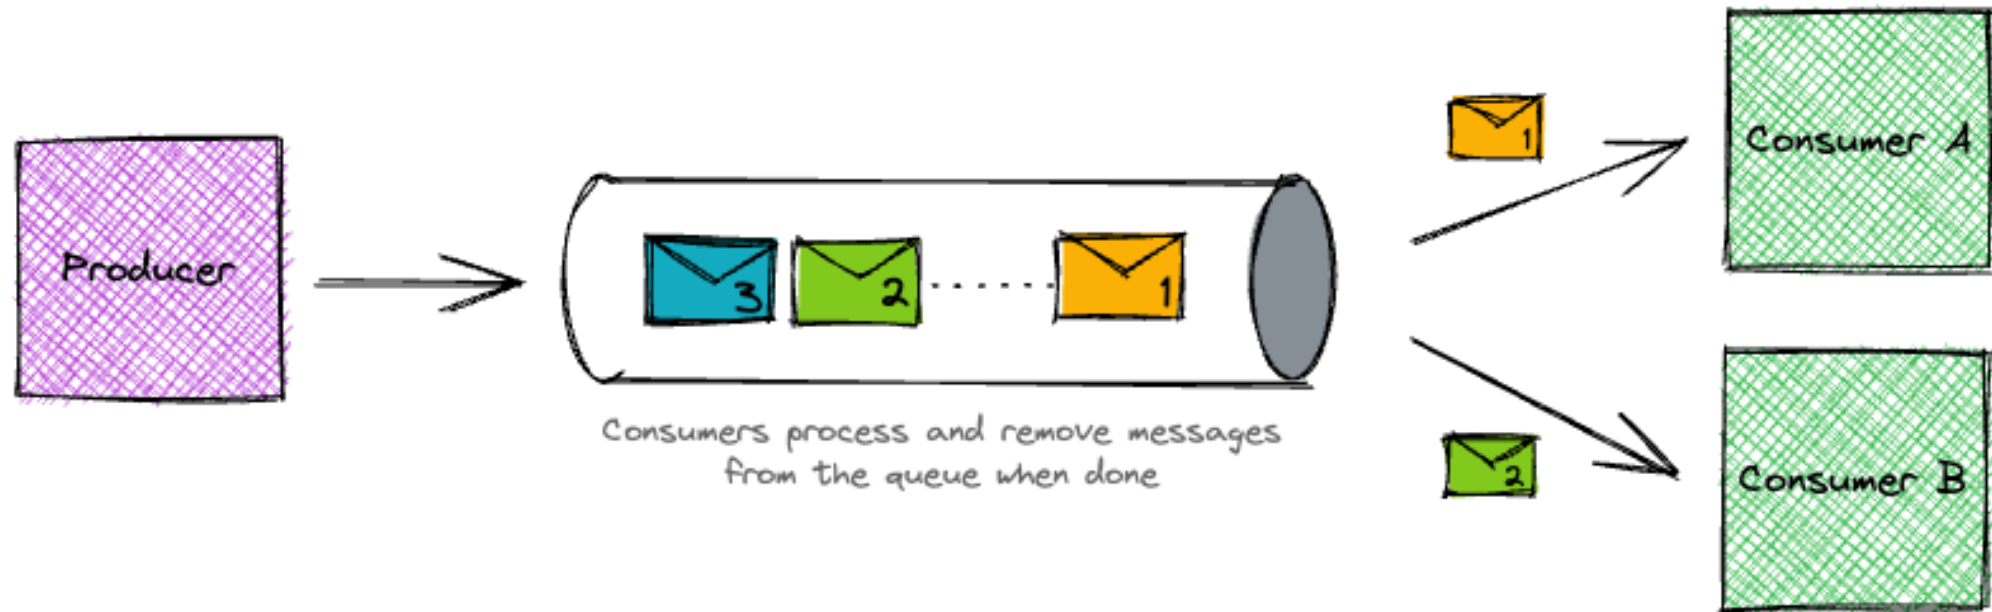


# Let's start with Why?





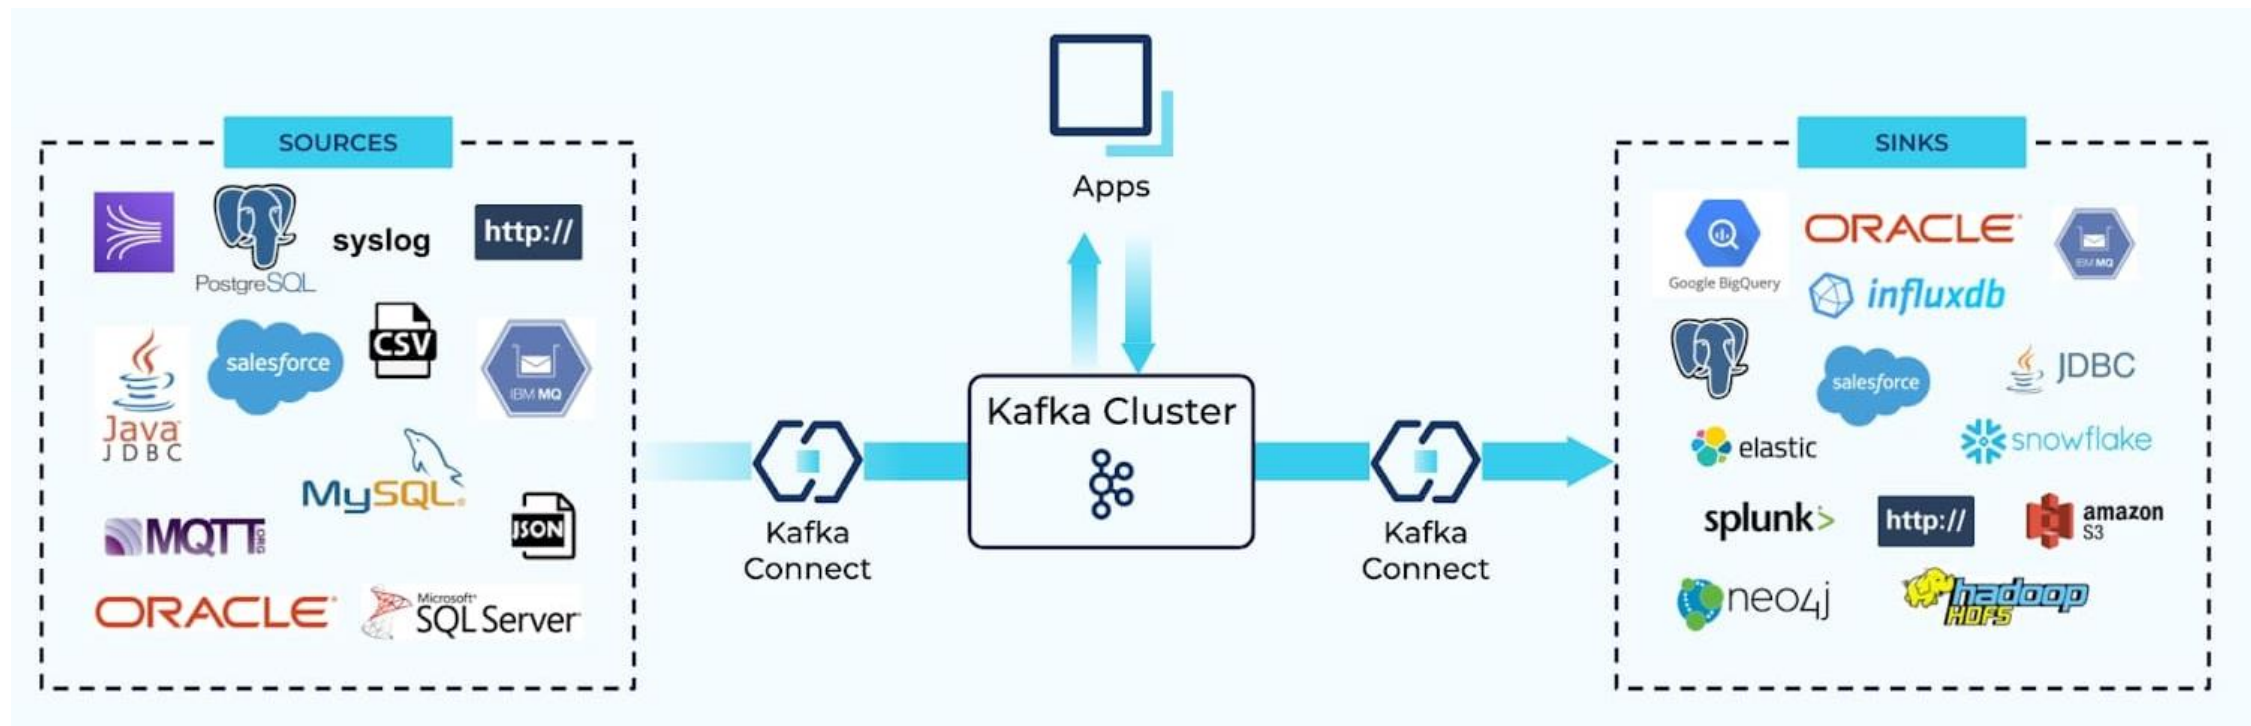


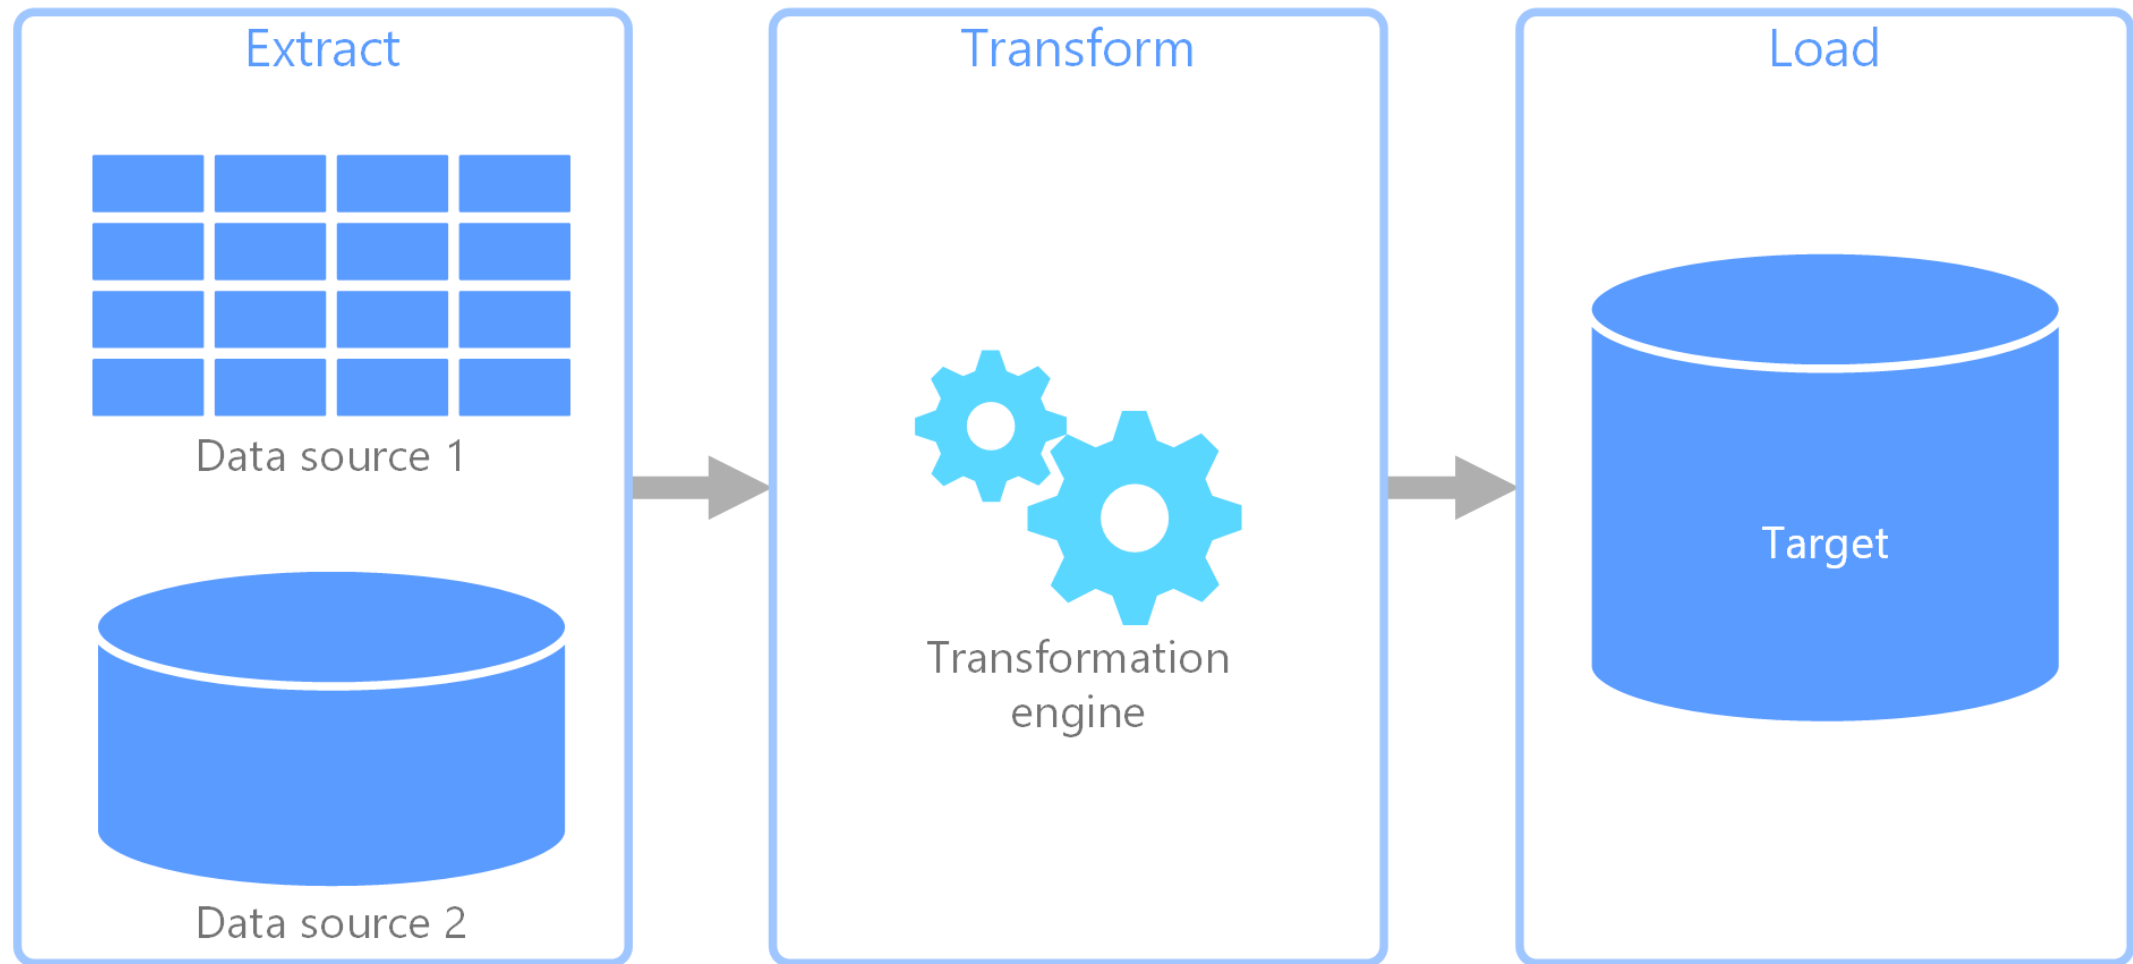


# Queues

Consumers pull messages, process and remove









# Overview

- Initially developed inside Linked-in
- Open-Source since 2011
- Donated to the Apache Software Foundation
- Event streaming platform
- Distributed
- Highly-available



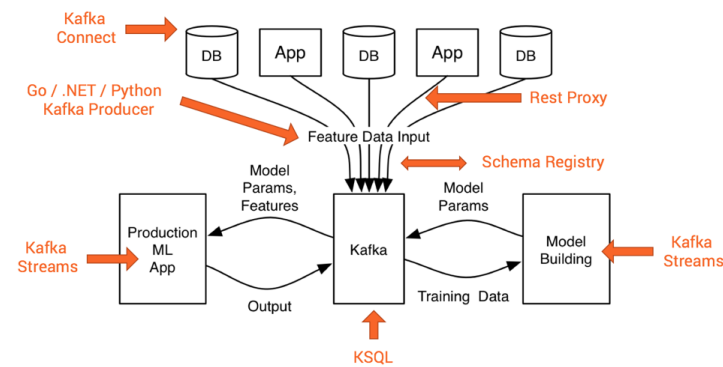
Real-time data  
processing




Stream  
processing



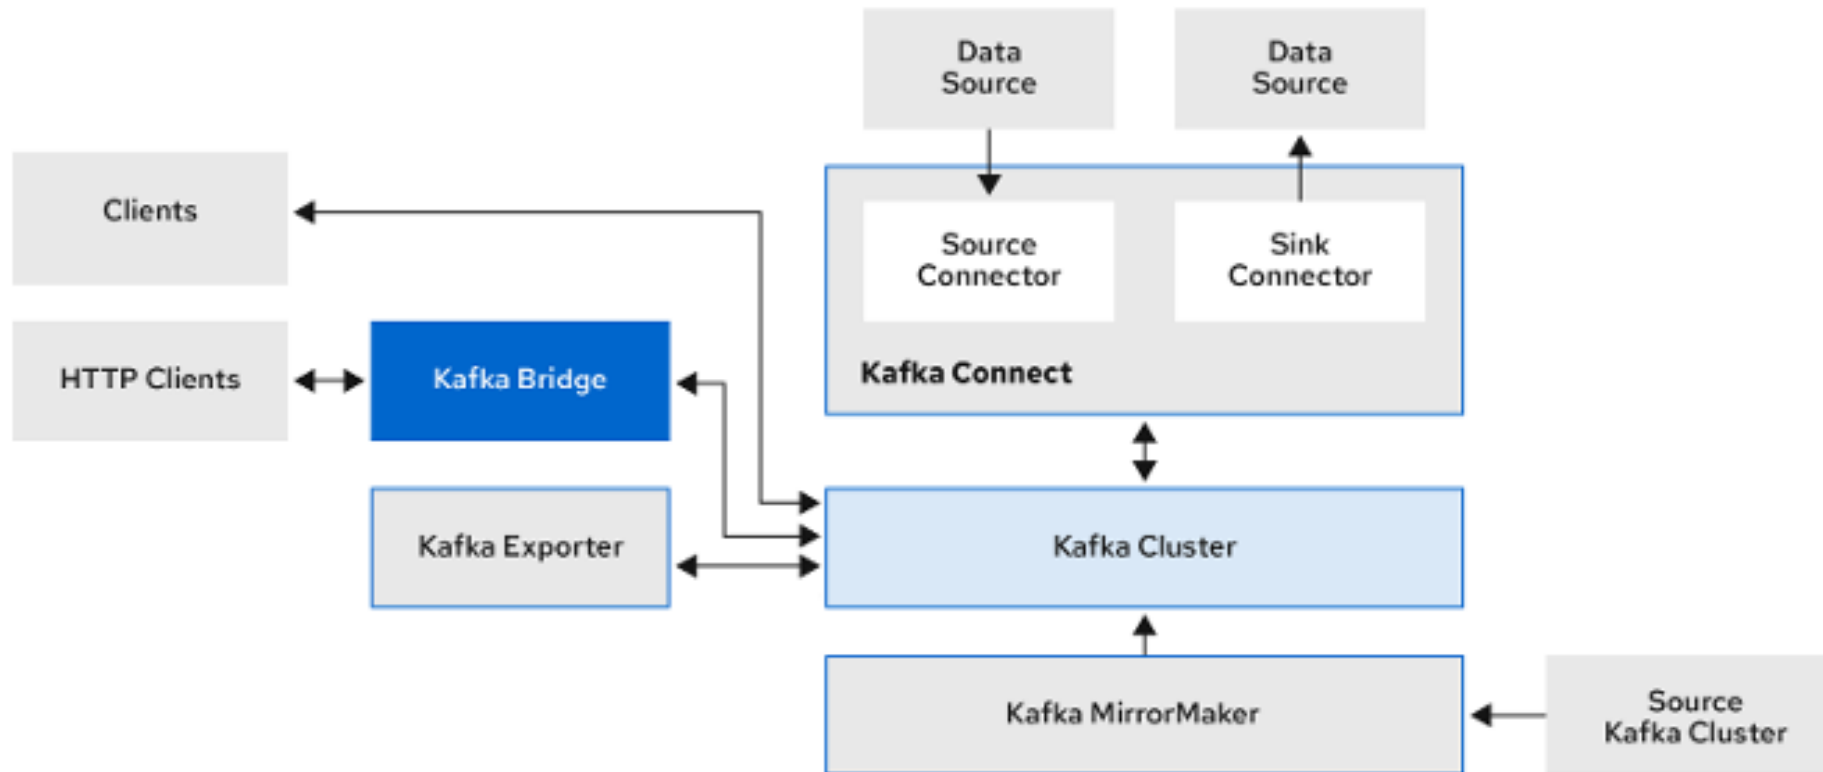
Building Data  
pipelines




Uber







Source: <https://strimzi.io/docs/operators/latest/overview>



Kafka Controller pool (Kraft)

## Kafka Cluster

Control Plane

Data Plane

**Controller**  
(active)

Metadata  
Cache

**Controller**

Metadata  
Cache

**Controller**

Metadata  
Cache

Broker

Metadata  
Cache

Broker

Metadata  
Cache

Broker

Metadata  
Cache

Broker

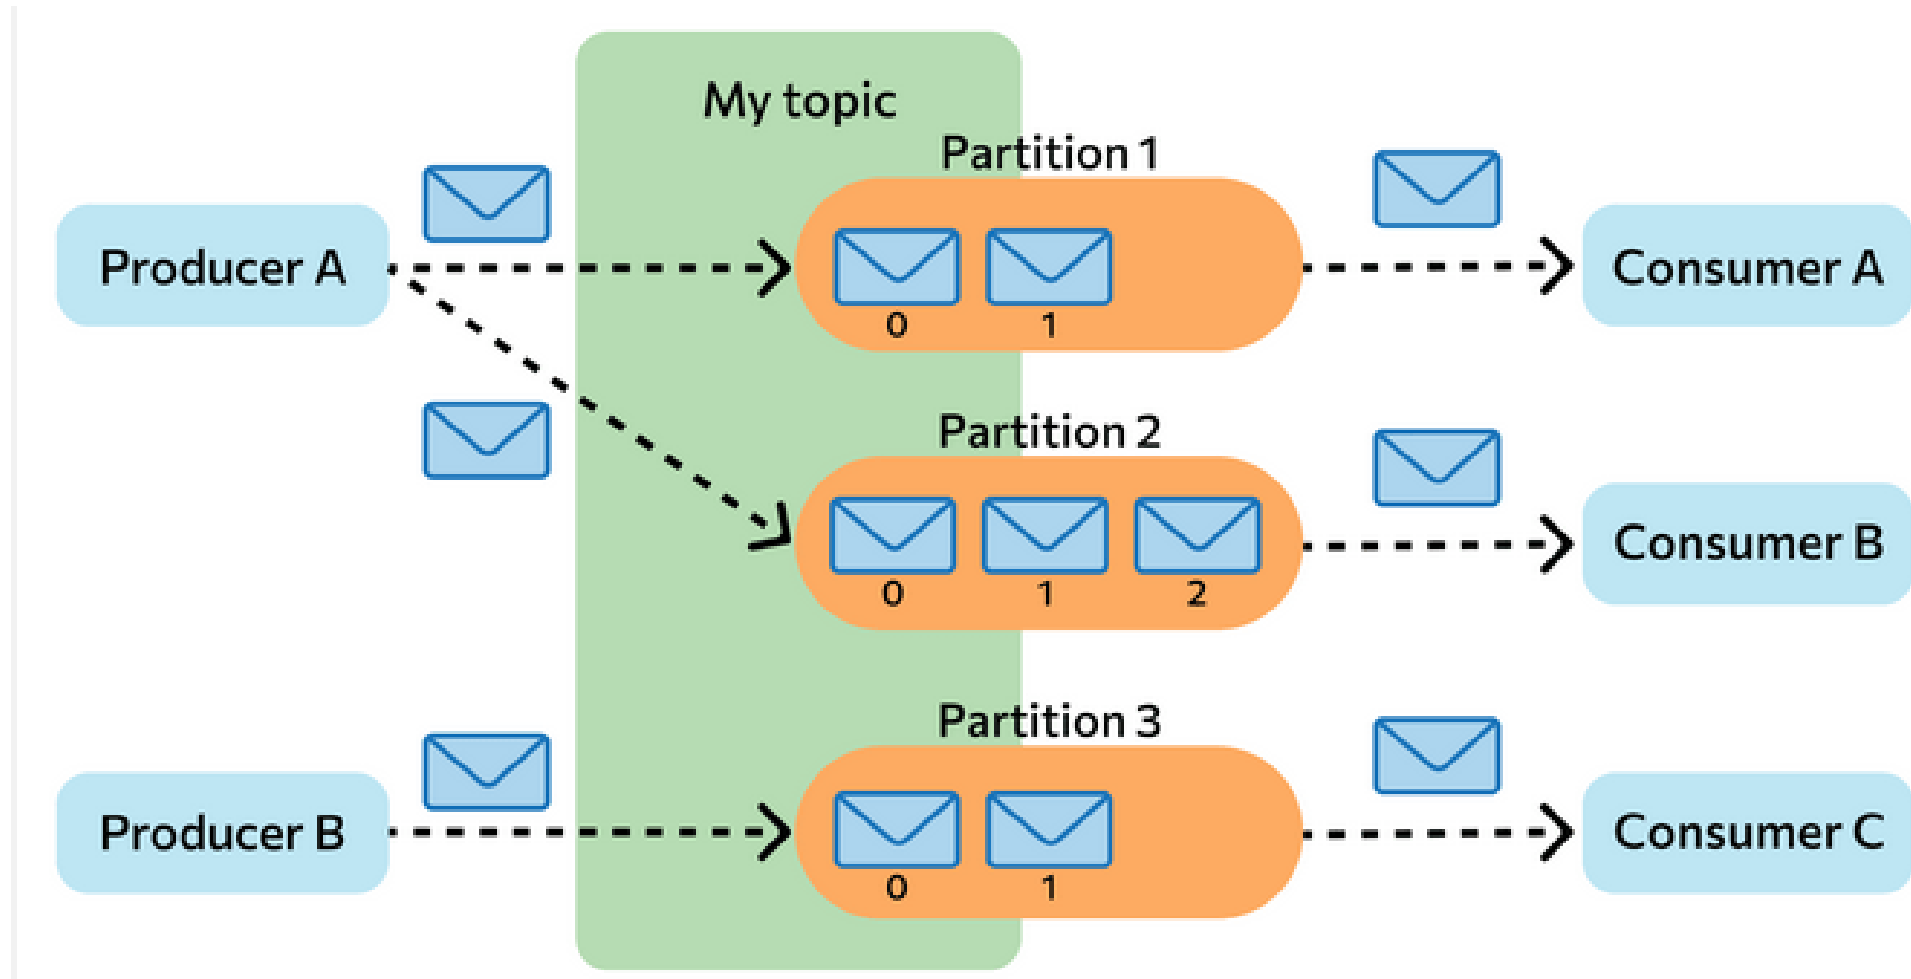
Metadata  
Cache

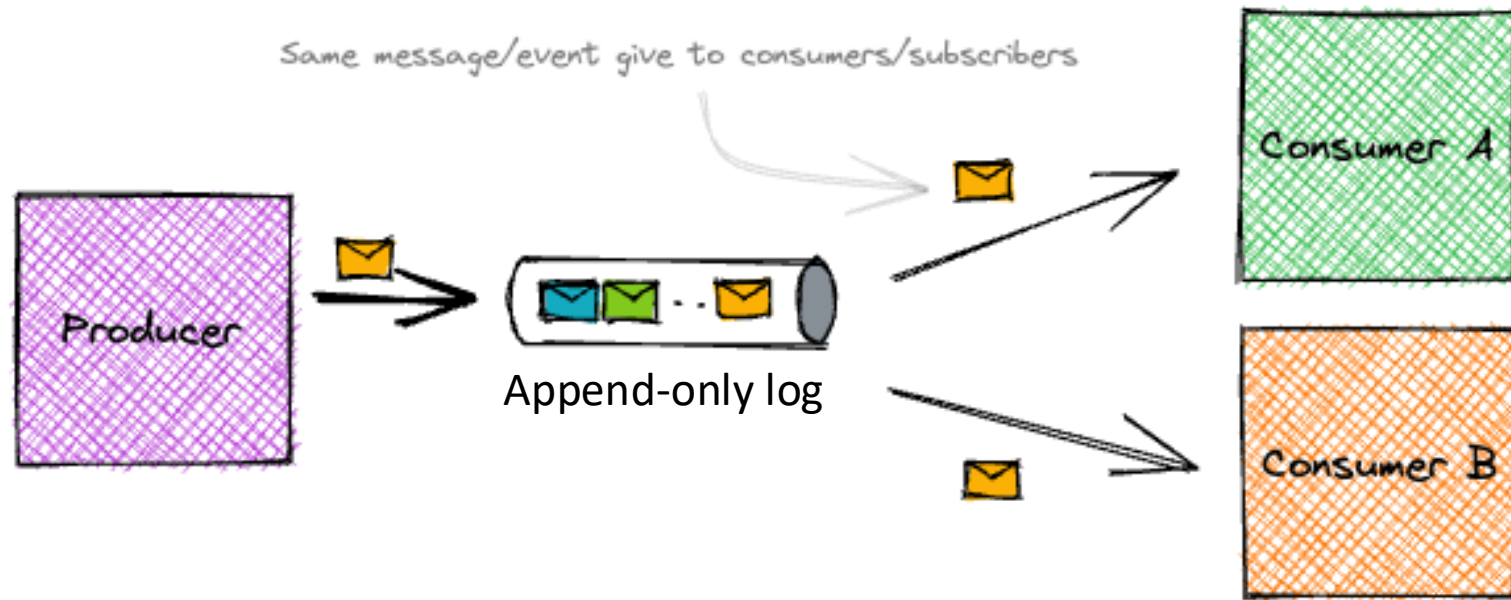
Broker

Metadata  
Cache

Broker

Metadata  
Cache

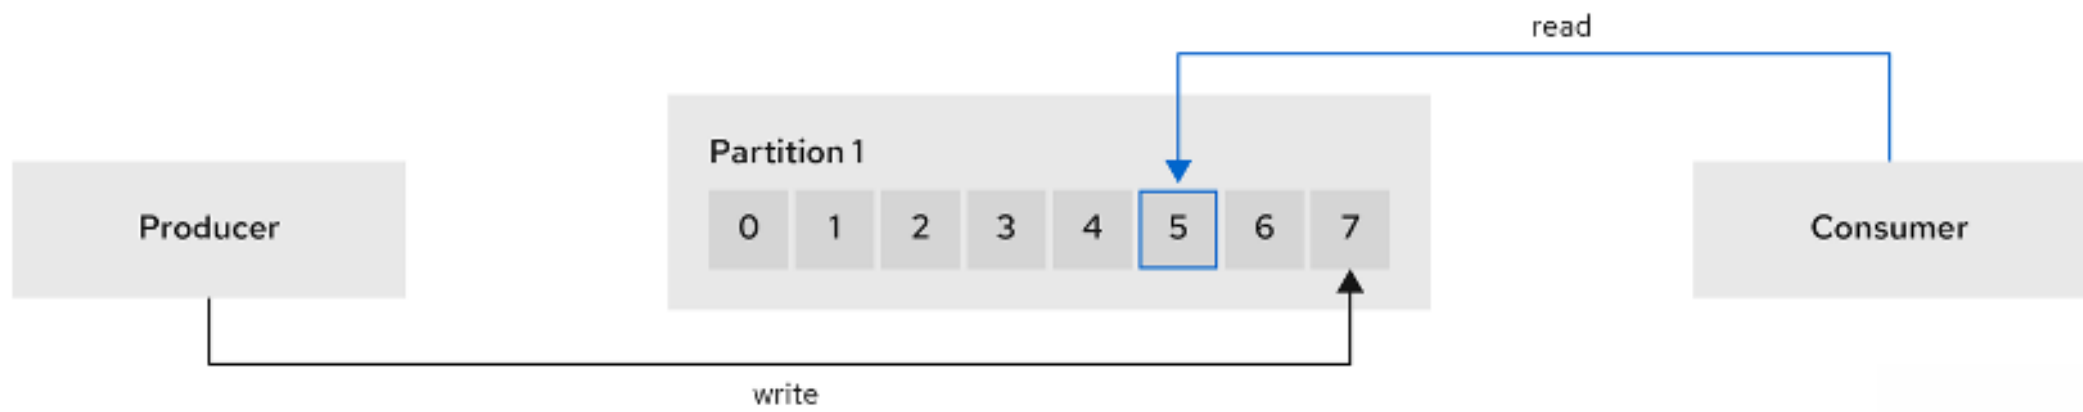





## Pub/Sub

Publish messages/events to many subscribers  
Each subscriber gets copy of event to process





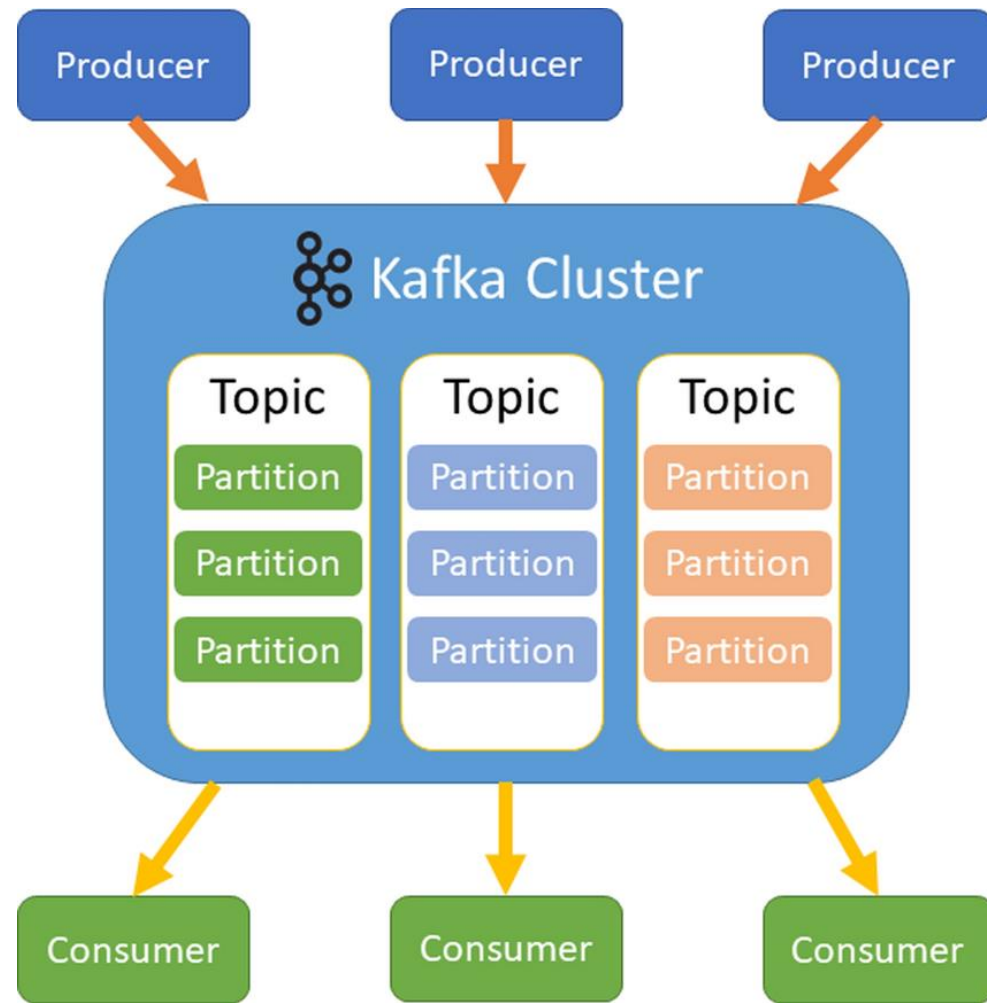
Consumers operate on  
stream of events

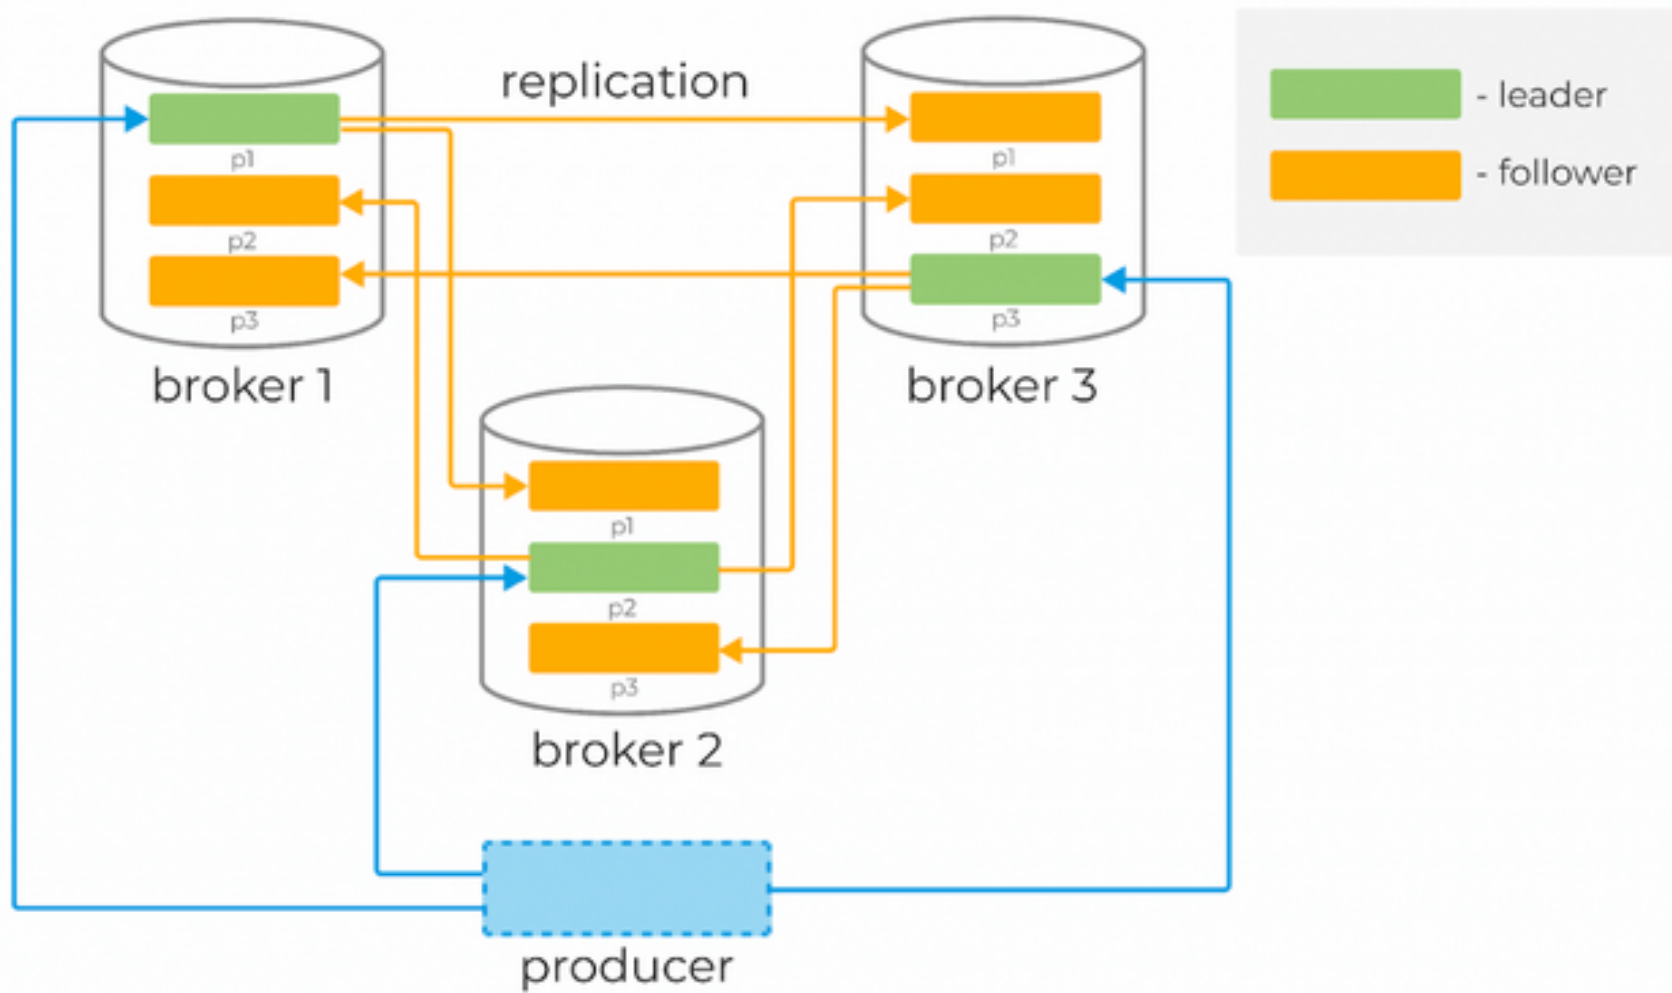


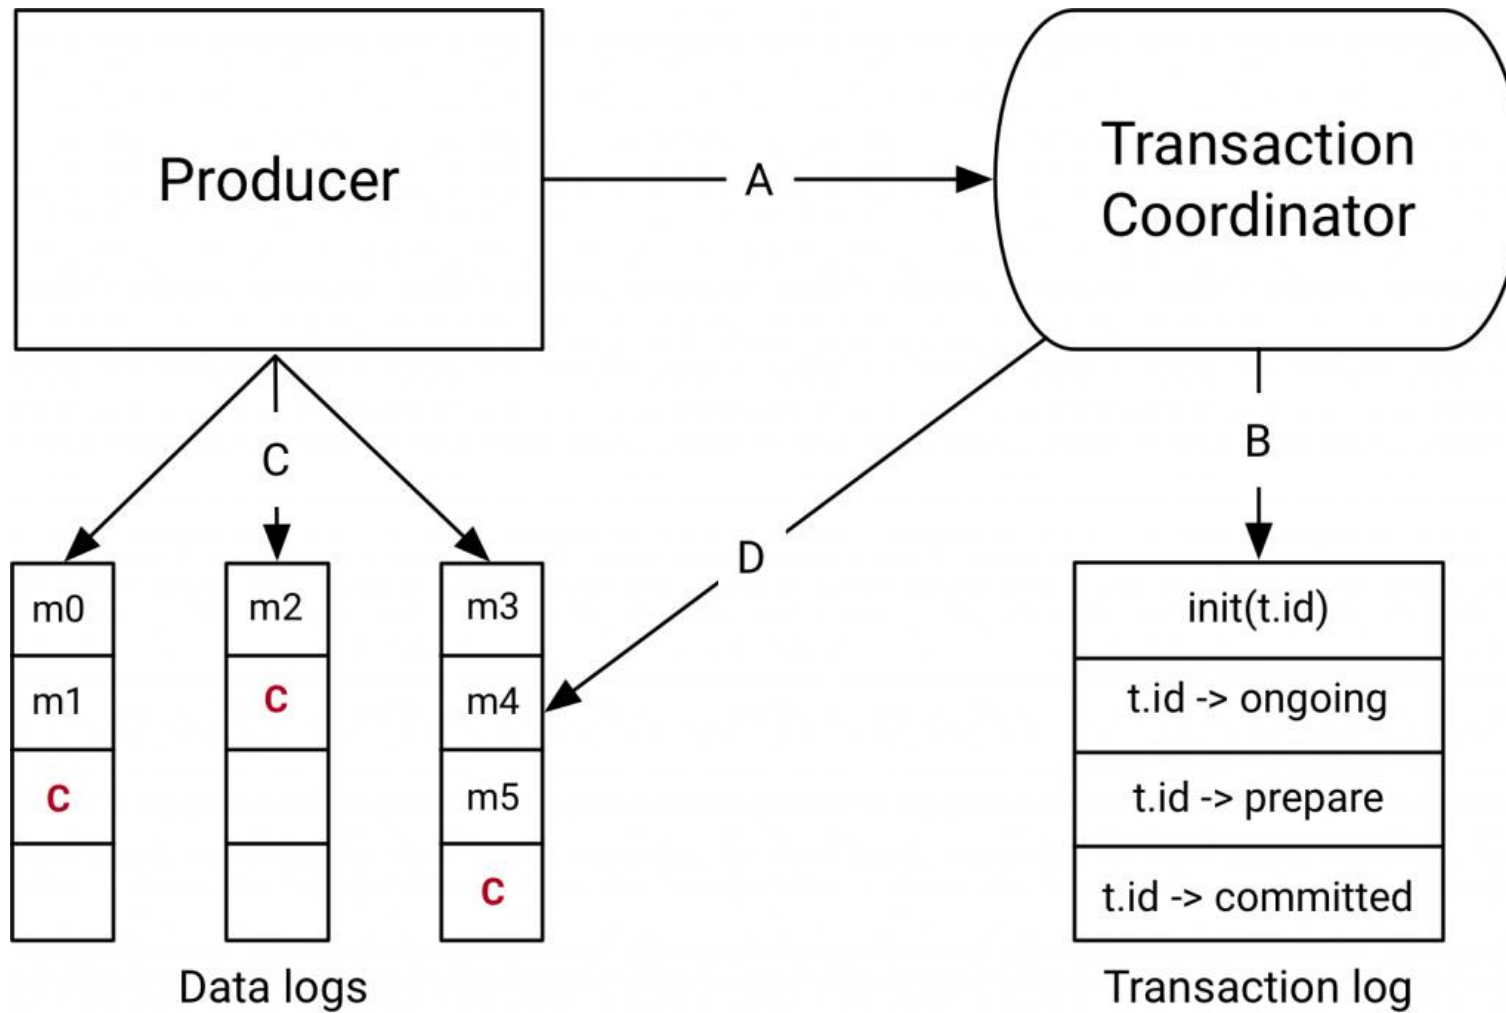
# Streams

Unbounded series of events

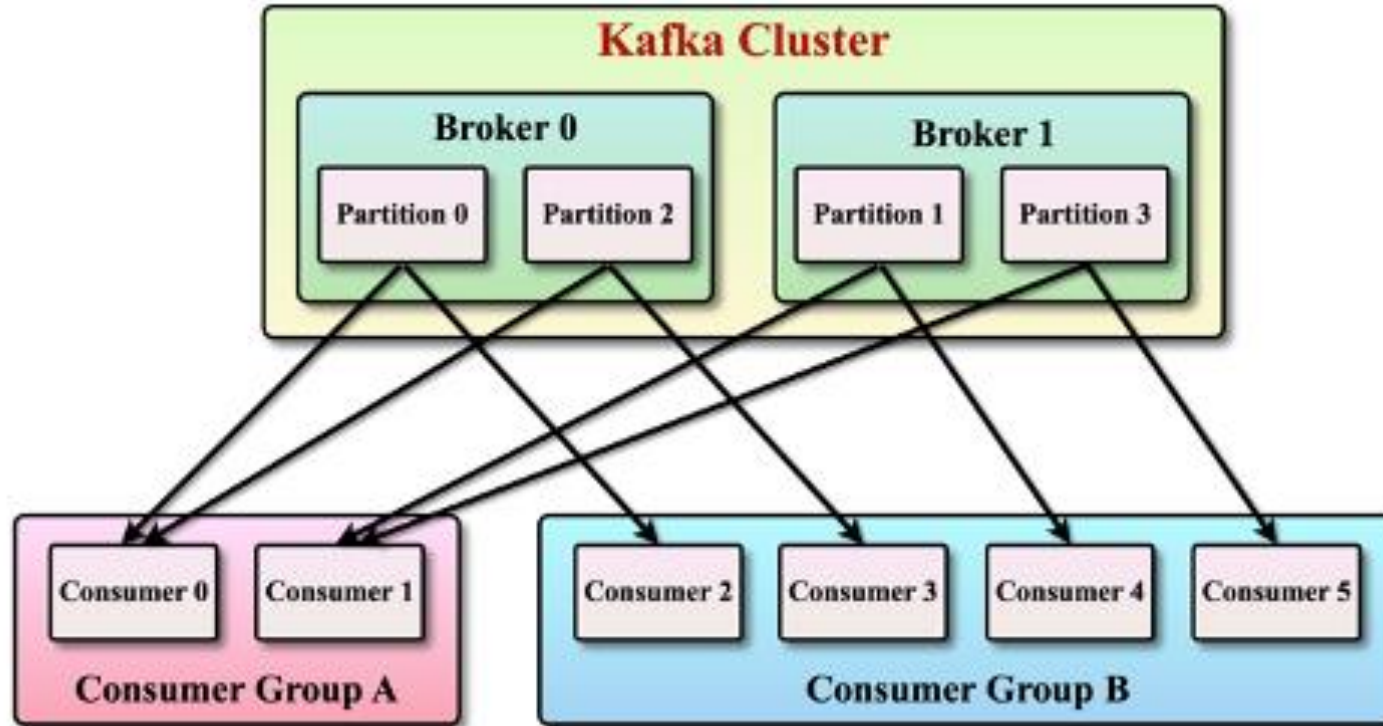
Common examples include user click streams / IOT / transactions

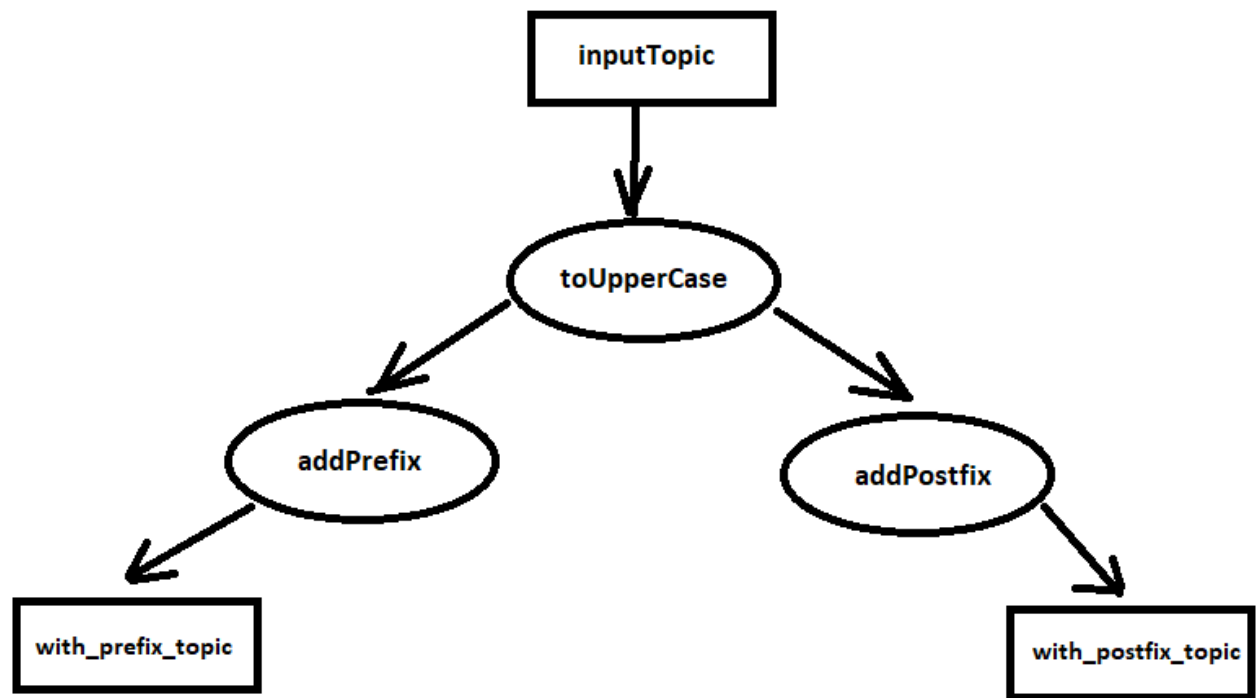















"Please process this at some point"



## Reducing knowledge of systems


Using messages/events allows us to be decoupled.  
This can reduce senders/producers info about the consumers.



## Parallel processing

With Pub/Sub notify downstream consumers  
Scale as downstream consumers process in parallel


## Problems that Kafka Message Brokers solve



## Reduce pressure off downstream consumers

Protect services getting overloaded downstream.  
Patterns to consume events at own ingestion rate.






## Prevent messages/data being lost

Messaging pattern is highly scalable  
Many consumers can process queue

# Apache Kafka on Kubernetes



# Running Apache Kafka on Kubernetes





## PROJECTS

# Strimzi

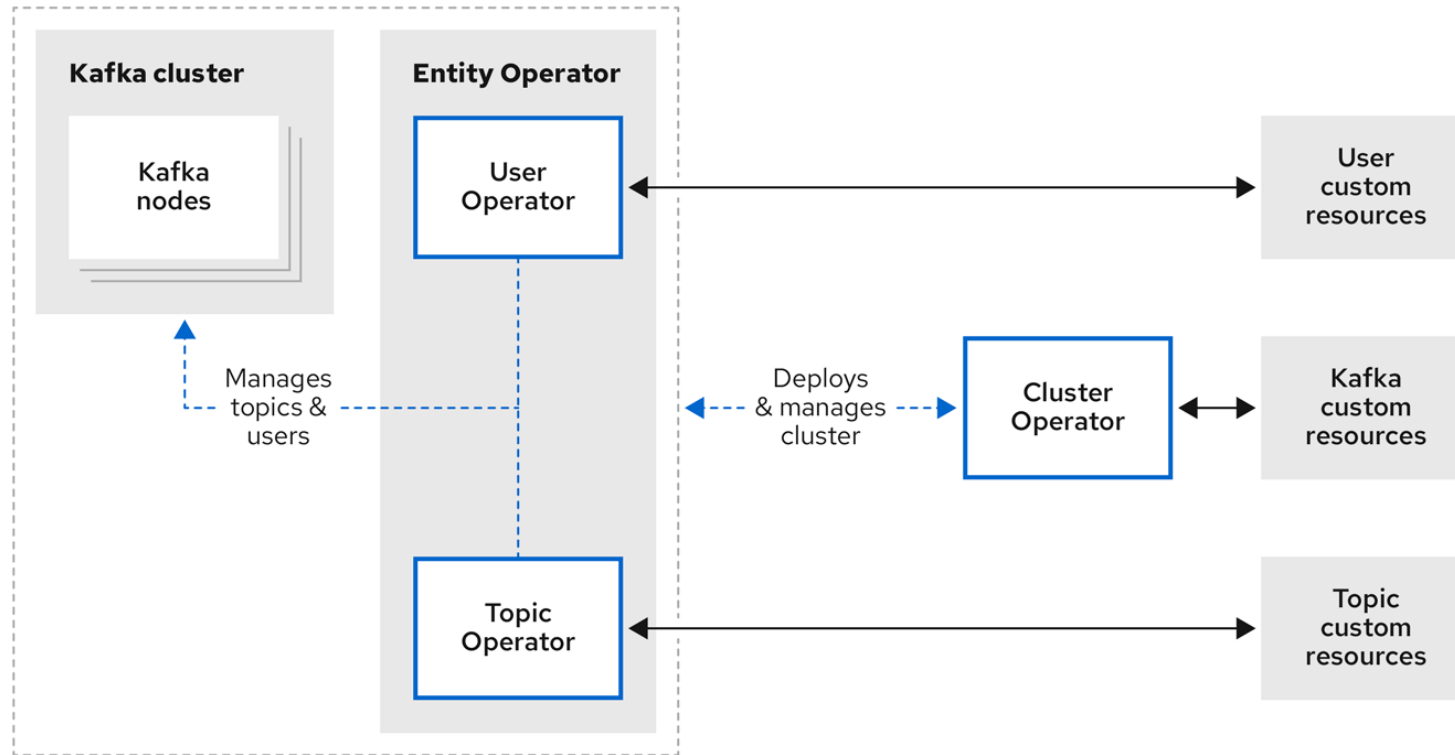


## Apache Kafka® running on Kubernetes

Strimzi was accepted to CNCF on August 28, 2019 and moved to the **Incubating** maturity level on February 8, 2024.

[VISIT PROJECT WEBSITE](#)

## Strimzi



## Kubernetes

574\_0424

# Strimzi Kafka Deployment Options

Installation method	Description
<a href="#">Deployment files (YAML files)</a>	<a href="#">Download the deployment files</a> to manually deploy Strimzi components. For the greatest flexibility, choose this method.
<a href="#">OperatorHub.io</a>	Deploy the Strimzi Cluster operator through the OperatorHub.io, then deploy Strimzi components using custom resources. This method provides a standard configuration and allows you to take advantage of automatic updates.
<a href="#">Helm chart</a>	Use a Helm chart to deploy the Cluster Operator, then deploy Strimzi components using custom resources. Helm charts provide a convenient way to manage the installation of applications.

```
$ helm repo add strimzi https://strimzi.io/charts  
$ helm repo update
```

```
$ helm install strimzi-kafka-operator strimzi/strimzi-kafka-operator \  
--version 0.45.0 -n kafka --create-namespace \  
-f ./strimzi-operator/strimzi-kafka-values.yaml
```

```
watchAnyNamespace: true
```

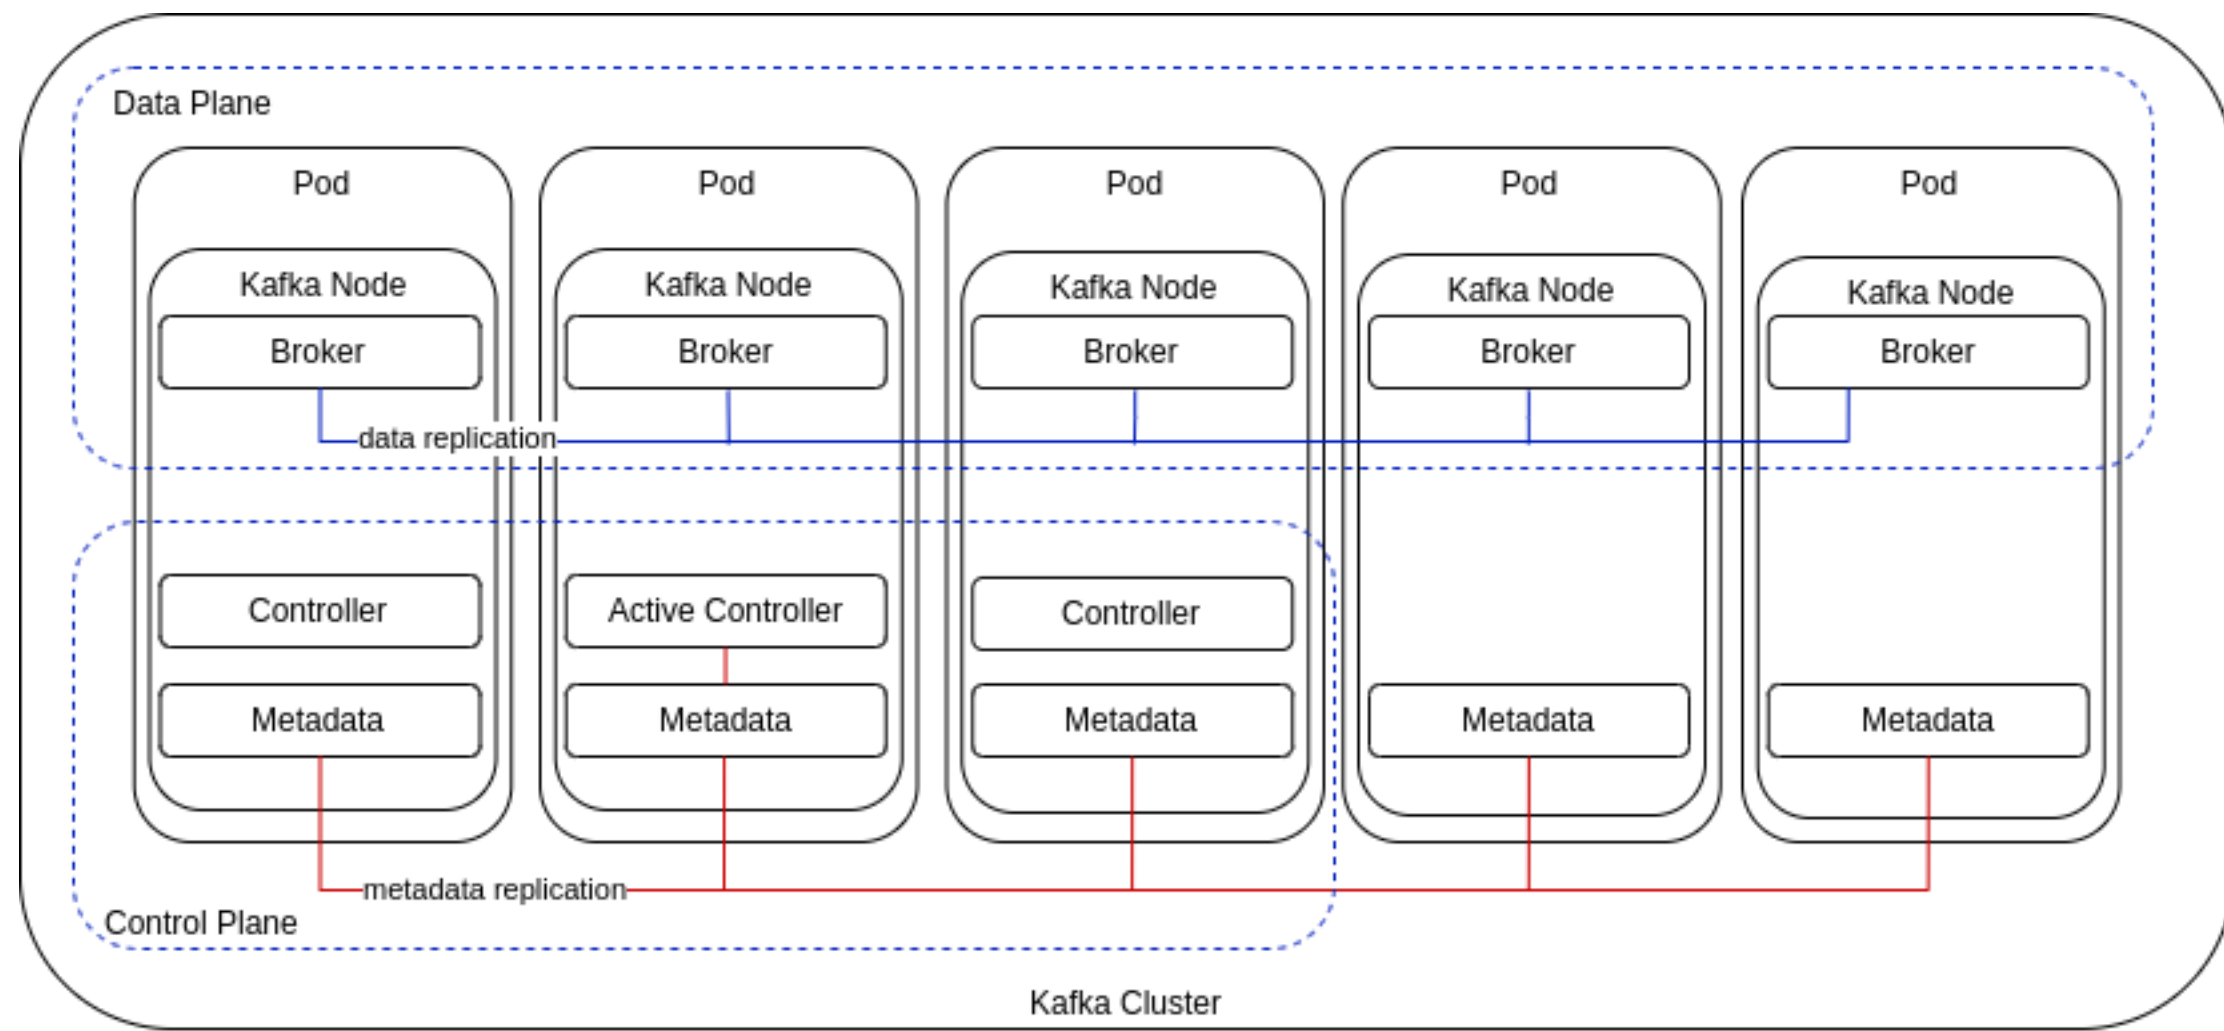
```
dashboards:
```

```
  enabled: true
```

```
  namespace: monitoring
```

Custom Helm Values





```
kubectl api-resources | grep strimzi
```

strimzipodsets	sps	core.strimzi.io/v1beta2	true	StrimziPodSet
kafkabridges	kb	kafka.strimzi.io/v1beta2	true	KafkaBridge
kafkaconnectors	kctr	kafka.strimzi.io/v1beta2	true	KafkaConnector
kafkaconnects	kc	kafka.strimzi.io/v1beta2	true	KafkaConnect
kafkamirrormaker2s	kmm2	kafka.strimzi.io/v1beta2	true	KafkaMirrorMaker2
kafkamirrormakers	kmm	kafka.strimzi.io/v1beta2	true	KafkaMirrorMaker
kafkanodepools	knp	kafka.strimzi.io/v1beta2	true	KafkaNodePool
kafkarebalances	kr	kafka.strimzi.io/v1beta2	true	KafkaRebalance
kafkas	k	kafka.strimzi.io/v1beta2	true	Kafka
kafkatopics	kt	kafka.strimzi.io/v1beta2	true	KafkaTopic
kafkausers	ku	kafka.strimzi.io/v1beta2	true	KafkaUser

```
apiVersion: kafka.strimzi.io/v1beta2
kind: Kafka
metadata:
  name: meet-kafka-cluster
  namespace: meet-kafka-cluster
  annotations:
    strimzi.io/node-pools: enabled
    strimzi.io/kraft: enabled
```

```
metricsConfig:
  type: jmxPrometheusExporter
  valueFrom:
    configMapKeyRef:
      name: kafka-metrics
      key: kafka-metrics-config.yml
entityOperator:
  topicOperator: {}
  userOperator: {}
  cruiseControl: {}
```

```
spec:
  kafka:
    config:
      offsets.topic.replication.factor: 3
      transaction.state.log.replication.factor: 3
      transaction.state.log.min.isr: 2
      default.replication.factor: 3
      min.insync.replicas: 2
    listeners:
      - name: plain
        port: 9092
        type: internal
        tls: false
      - name: tls
        port: 9093
        type: internal
        tls: true
    version: 3.9.0
```

```
apiVersion: kafka.strimzi.io/v1beta2
kind: KafkaNodePool
metadata:
  name: xs-nodepool
  namespace: meet-kafka-cluster
  labels:
    strimzi.io/cluster: meet-kafka-cluster
spec:
  replicas: 3
  roles:
    - controller
    - broker
  storage:
    type: jbod
    volumes:
      - id: 0
        type: persistent-claim
        size: 1Gi
        deleteClaim: false
```

Persistent storage types:

- **persistent-claim** - for a single persistent volume
- **Jbod** - for multiple persistent volumes in a Kafka cluster

Ephemeral storage types:

- **ephemeral** - uses emptyDir Volumes

Kafka Node Roles:

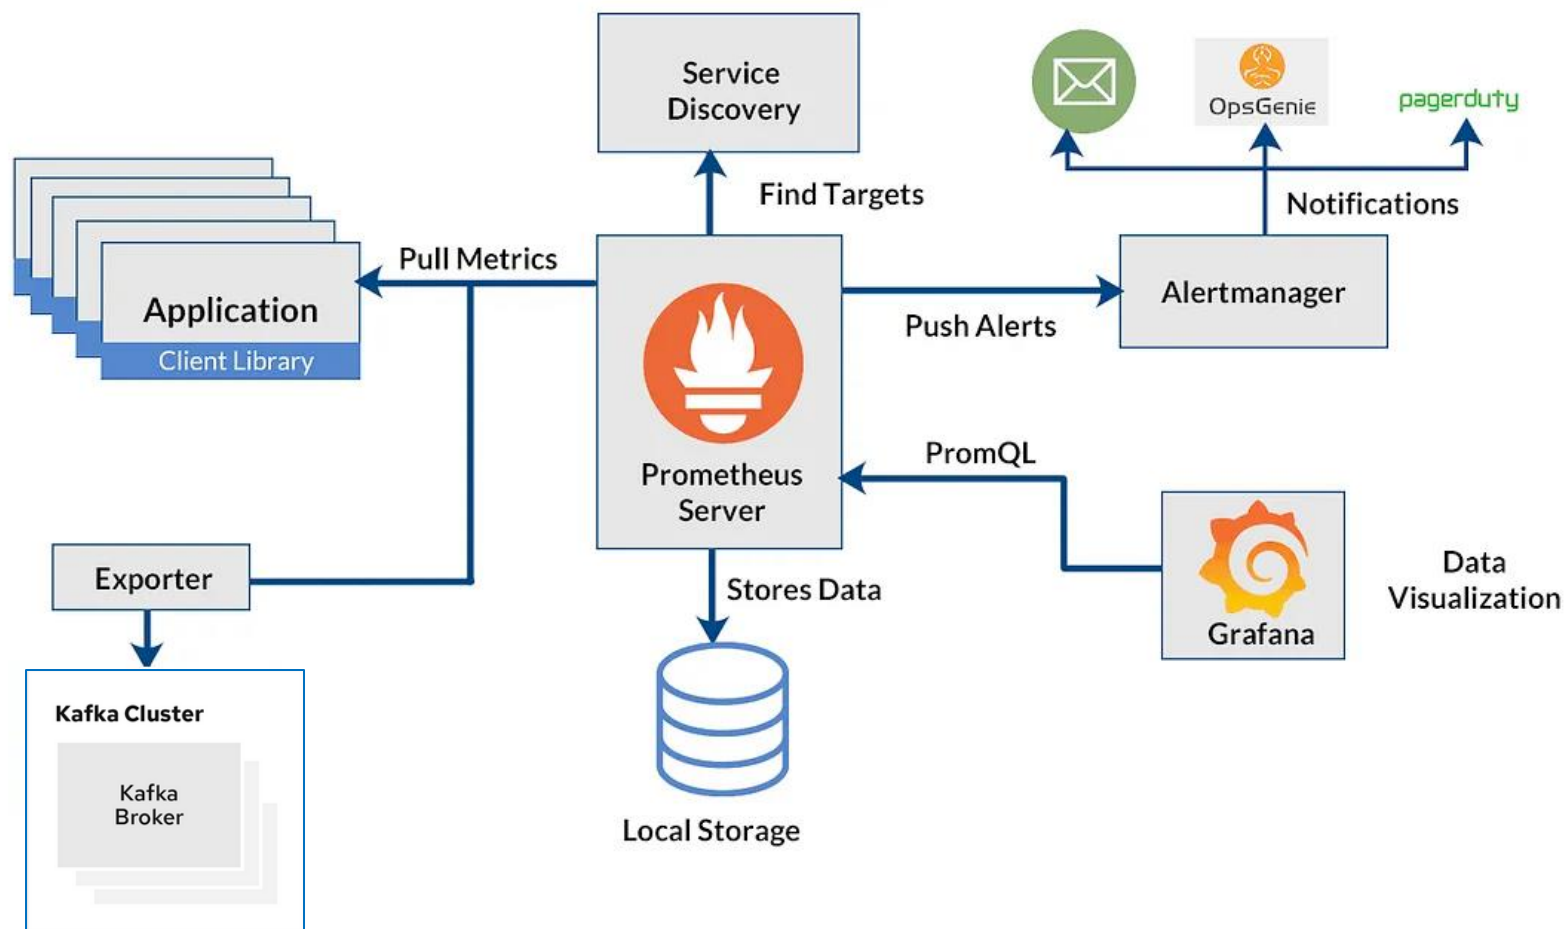
- **Controller** - Manage cluster metadata and state
- **Broker** - Manage message data
- **Dual-Role** - All-in-one

<https://strimzi.io/docs/operators/latest/full/deploying.html#con-config-storage-kraft-str>

# Monitoring Apache Kafka on Kubernetes

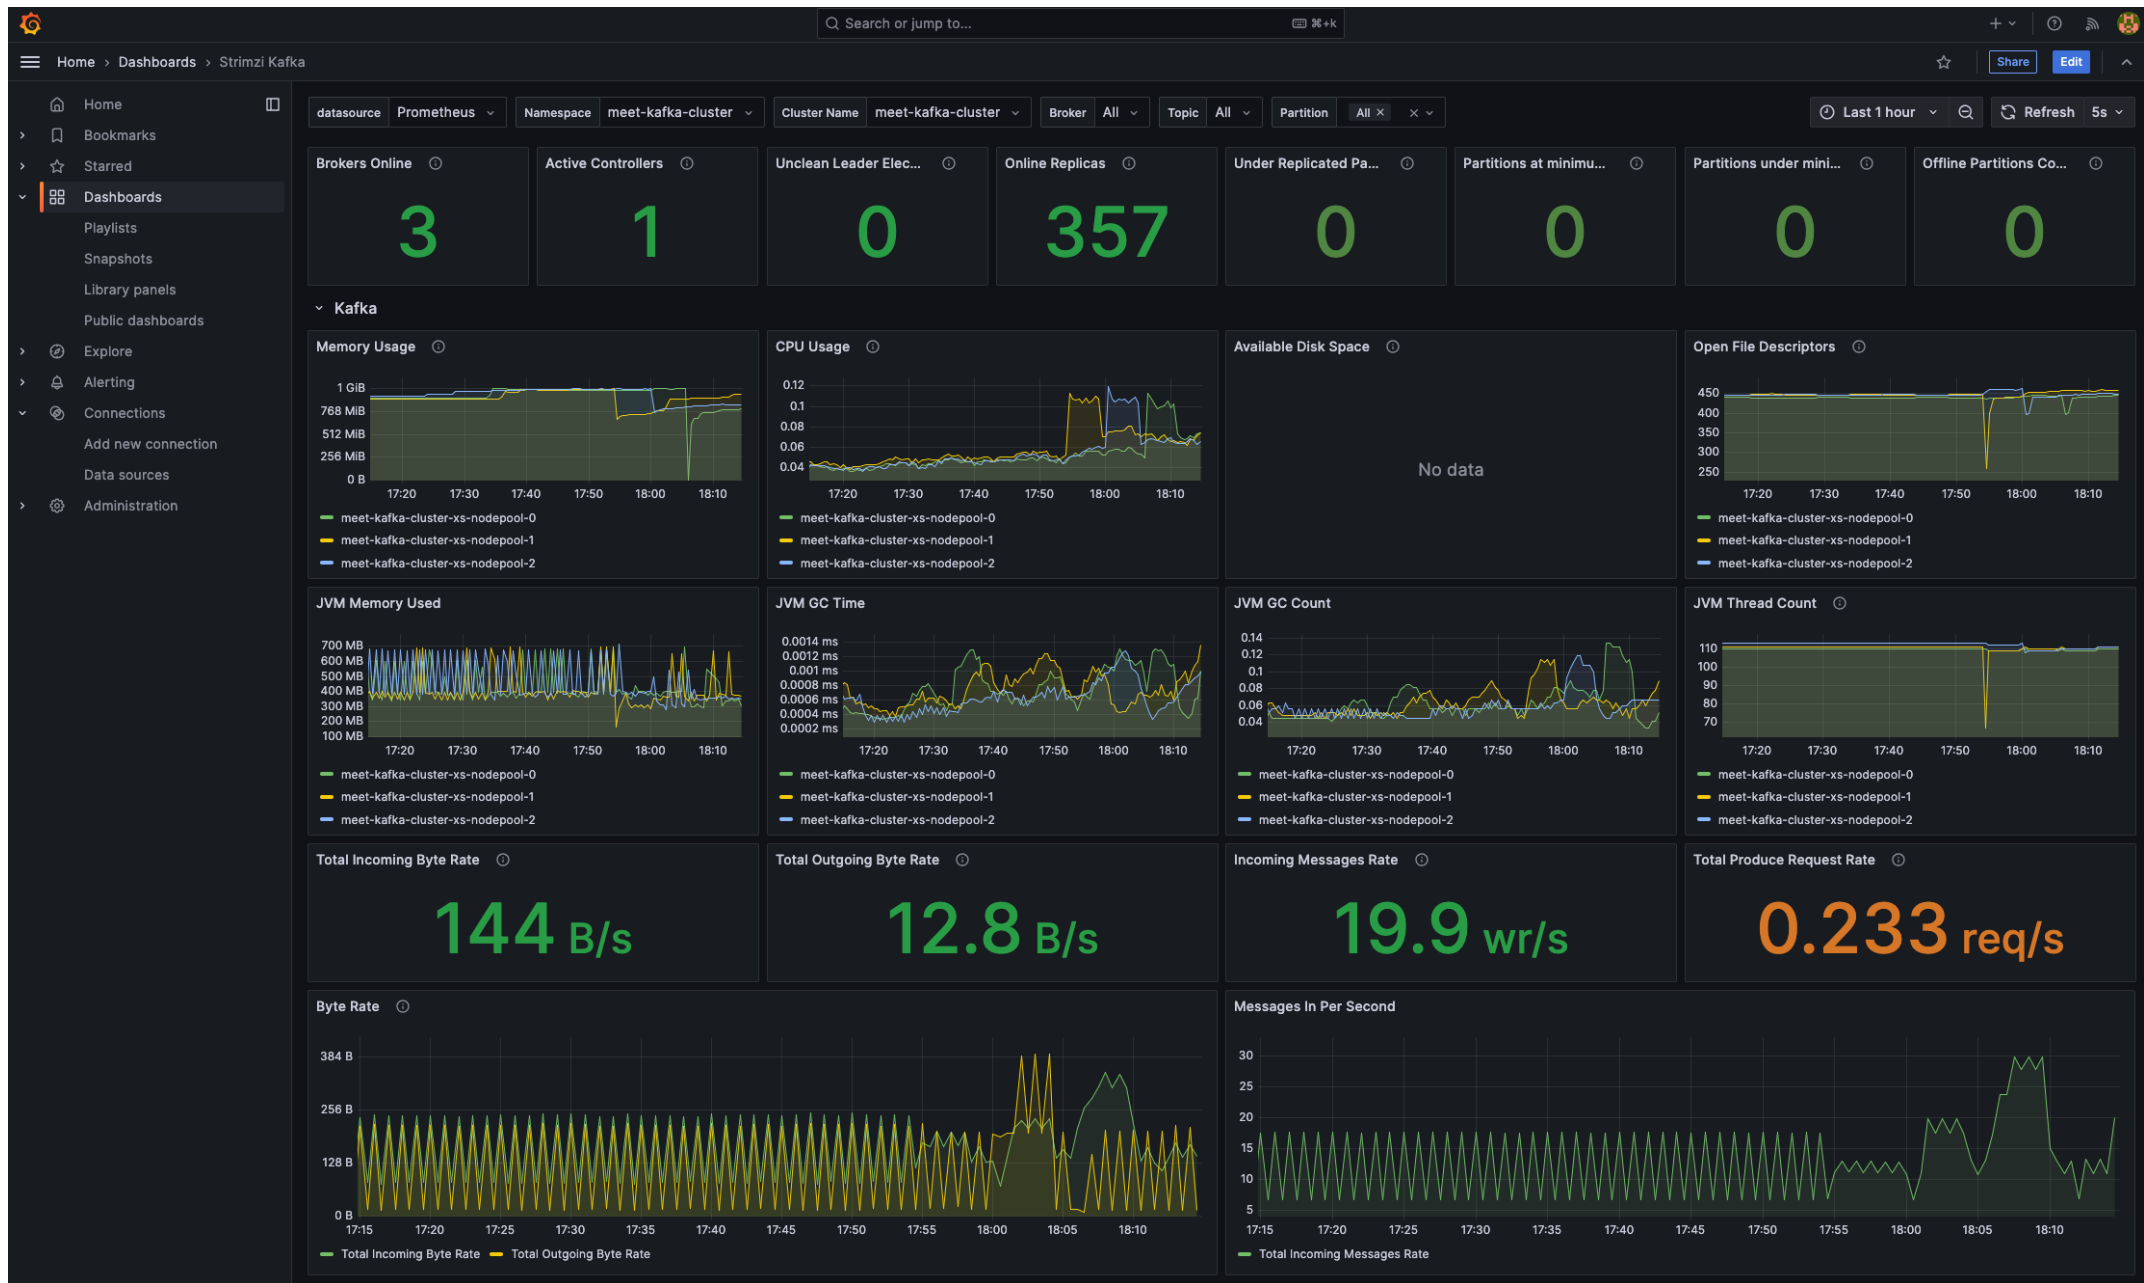


# Monitoring Kafka



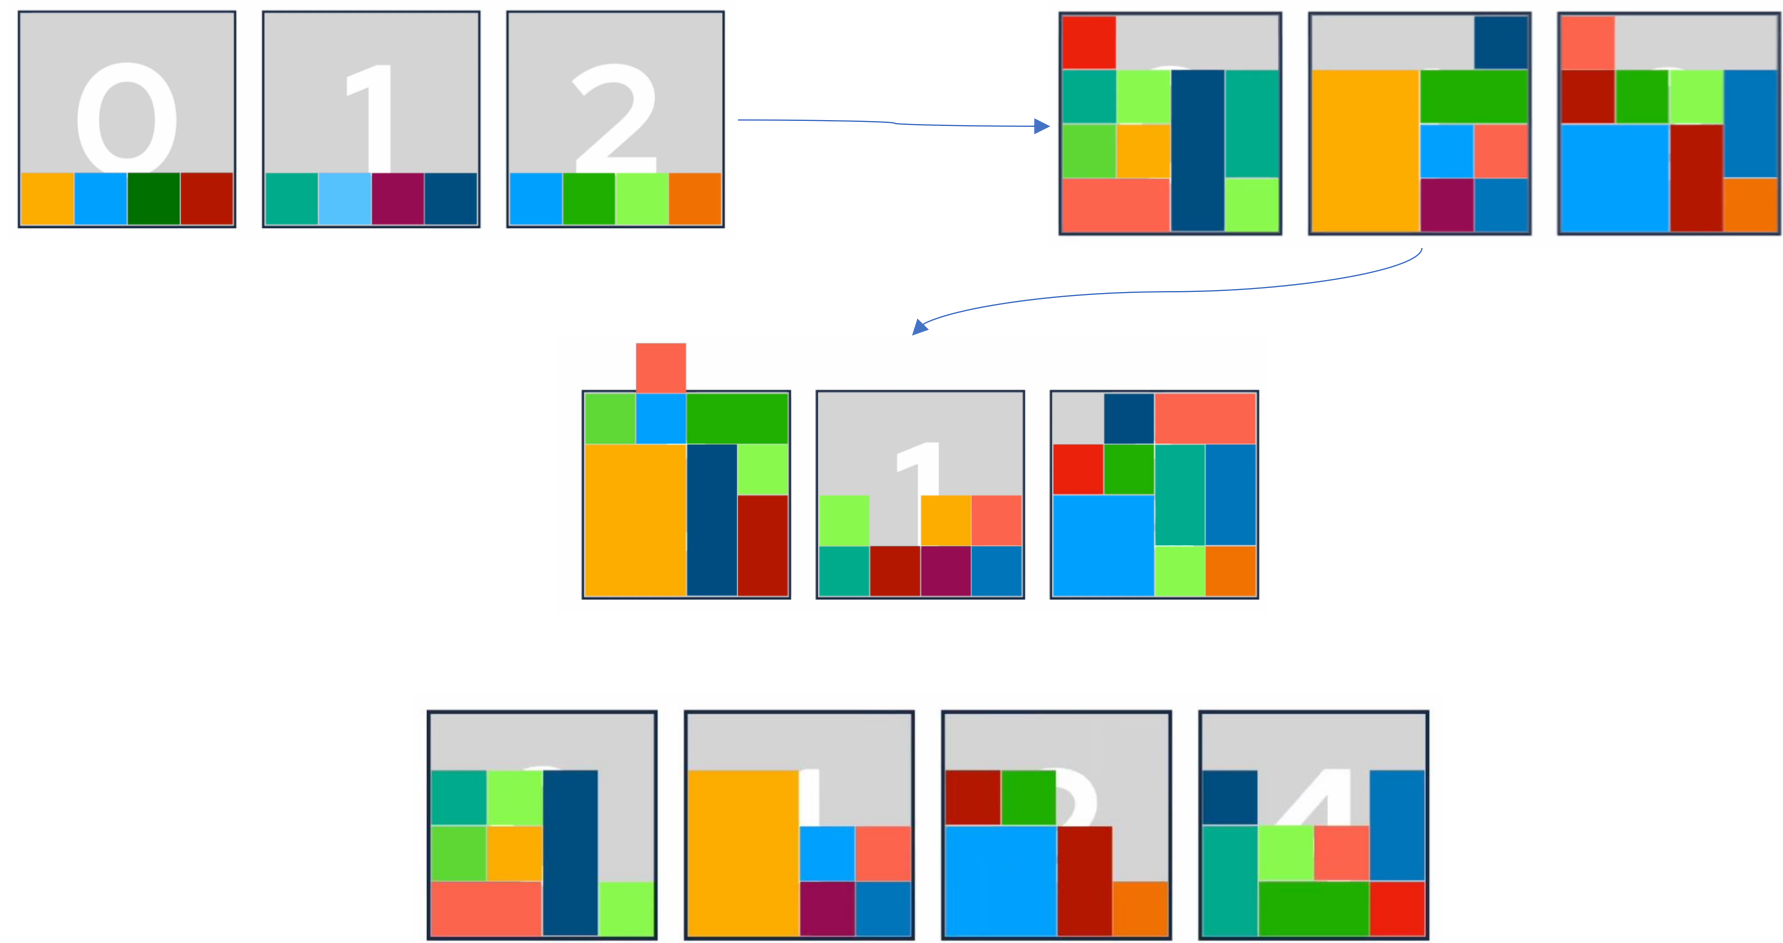
```
➤ kubectl -n monitoring get cm -l grafana_dashboard=1 |grep strimzi
```

strimzi-cruise-control	1	3d23h
strimzi-kafka	1	3d23h
strimzi-kafka-bridge	1	3d23h
strimzi-kafka-connect	1	3d23h
strimzi-kafka-exporter	1	3d23h
strimzi-kafka-mirror-maker-2	1	3d23h
strimzi-kafka-oauth	1	3d23h
strimzi-kraft	1	3d23h
strimzi-operators	1	3d23h
strimzi-zookeeper		





# Broker balancing



# Kafka Cluster Balancing

- What is the Optimal partitions/replica Distribution?
  - Rack Awareness?
  - Distribution of Leader replicas
  - Optimal Network, CPU, storage, or RAM utilization?
- Manual or automated optimizations?

```
apiVersion: kafka.strimzi.io/v1beta2
kind: KafkaRebalance
metadata:
  name: meet-cluster-rebalance
  labels:
    strimzi.io/cluster: meet-kafka-cluster
# Use the default Cruise Control optimization goals
spec: {}
```

## CruiseControl CRD

---

NAME	CLUSTER	TEMPLATE	STATUS
meet-cluster-rebalance	meet-kafka-cluster		ProposalReady

Name: meet-cluster-rebalance  
Namespace: meet-kafka-cluster  
Labels: strimzi.io/cluster=meet-kafka-cluster  
Annotations: <none>  
API Version: kafka.strimzi.io/v1beta2  
Kind: KafkaRebalance  
Metadata:  
Creation Timestamp: 2025-01-20T21:30:18Z  
Generation: 1  
Resource Version: 580148  
UID: d836fa04-4715-4f48-b68c-84dd0b6dff37  
Spec:  
Status:  
Conditions:  
Last Transition Time: 2025-01-20T21:30:18.754838993Z  
Status: True  
Type: ProposalReady

Observed Generation: 1  
Optimization Result:  
After Before Load Config Map: meet-cluster-rebalance  
Data To Move MB: 0  
Excluded Brokers For Leadership:  
Excluded Brokers For Replica Move:  
Excluded Topics:  
Intra Broker Data To Move MB: 0  
Monitored Partitions Percentage: 100  
Num Intra Broker Replica Movements: 0  
Num Leader Movements: 3  
Num Replica Movements: 0  
On Demand Balancedness Score After: 89.4347095948149  
On Demand Balancedness Score Before: 89.4347095948149  
Provision Recommendation:  
Provision Status: RIGHT\_SIZED  
Recent Windows: 1  
...


# Best Practices

- Install the Strimzi operator in a separate namespace from the Kafka cluster and other Kafka components it manages
  - Ensure clear separation of resources and configurations.
  - Avoid the issues associated with installing multiple Strimzi operators in a Kubernetes cluster
- Use a single Strimzi operator to manage all your Kafka instances within a Kubernetes cluster.
- Update the Strimzi operator and the supported Kafka version as often as possible
  - to reflect the latest features and enhancements.

# Best Practices

- By default, a single replica of the Cluster Operator is deployed.
  - Add extra stand-by replicas in case of disruption.
  - One replica is elected leader
- The Cluster Operator watches for updates in the namespaces where the Kafka resources are deployed.
  - We can specify which namespaces to watch in the Kubernetes cluster.
  - Watching multiple selected namespaces has the most impact on performance due to increased processing overhead.
- To optimize performance for namespace monitoring, it is generally recommended to either watch:
  - a single namespace
  - monitor the entire cluster.

# Thank you!



<https://dojobits.io/>